

# Apache Spark for library developers

William Benton  
willb@redhat.com  
@willb

Erik Erlandson  
eje@redhat.com  
@manyangled


# The Silex and Isarn libraries

Reusable open-source code that works with Spark, factored from internal apps.



We've tracked Spark releases since Spark 1.3.0.

See <https://silex.radanalytics.io> and <http://isarnproject.org>









# Forecast

Basic considerations for reusable Spark code

Generic functions for parallel collections


Extending data frames with custom aggregates

Exposing JVM libraries to Python


Sharing your work with the world




# Basic considerations





























# Today's main themes



# Cross-building for Scala

in your SBT build definition:

```
scalaVersion := "2.11.11"
```

```
crossScalaVersions := Seq("2.10.6", "2.11.11")
```

in your shell:

```
$ sbt +compile
```

```
$ sbt "++ 2.11.11" compile
```

# Cross-building for Scala

in your SBT build definition:

```
scalaVersion := "2.11.11"
```

```
crossScalaVersions := Seq("2.10.6", "2.11.11")
```

in your shell:

```
$ sbt +compile           # or test, package, publish, etc.
```

```
$ sbt "++ 2.11.11" compile
```



# Cross-building for Scala

in your SBT build definition:

```
scalaVersion := "2.11.11"
```

```
crossScalaVersions := Seq("2.10.6", "2.11.11")
```

in your shell:

```
$ sbt +compile # or test, package, publish, etc.
```

```
$ sbt "++ 2.11.11" compile
```

# “Bring-your-own Spark”

in your SBT build definition:

```
libraryDependencies += Seq(  
  "org.apache.spark" %% "spark-core" % "2.3.0" % Provided,  
  "org.apache.spark" %% "spark-sql" % "2.3.0" % Provided,  
  "org.apache.spark" %% "spark-mllib" % "2.3.0" % Provided,  
  "joda-time" % "joda-time" % "2.7",  
  "org.scalatest" %% "scalatest" % "2.2.4" % Test)
```


# “Bring-your-own Spark”

in your SBT build definition:

```
libraryDependencies += Seq(  
  "org.apache.spark" %% "spark-core" % "2.3.0" % Provided,  
  "org.apache.spark" %% "spark-sql" % "2.3.0" % Provided,  
  "org.apache.spark" %% "spark-mllib" % "2.3.0" % Provided  
  "joda-time" % "joda-time" % "2.7",  
  "org.scalatest" %% "scalatest" % "2.2.4" % Test)
```




# Taking care with resources



# Taking care with resources



# Taking care with resources



# Caching when necessary

```
def step(rdd: RDD[_]) = {  
  
    rdd.cache()  
    result = trainModel(rdd)  
  
    result  
}
```

# Caching when necessary

```
def step(rdd: RDD[_]) = {  
  
    rdd.cache()  
    result = trainModel(rdd)  
  
    result  
}
```

# Caching when necessary

```
def step(rdd: RDD[_]) = {  
  
    rdd.cache()  
    result = trainModel(rdd)  
  
    rdd.unpersist()  
  
    result  
}
```



# Caching when necessary

```
def step(rdd: RDD[_]) = {  
    val wasUncached = rdd.storageLevel == StorageLevel.NONE  
    if (wasUncached) { rdd.cache() }  
    result = trainModel(rdd)  
  
    result  
}
```

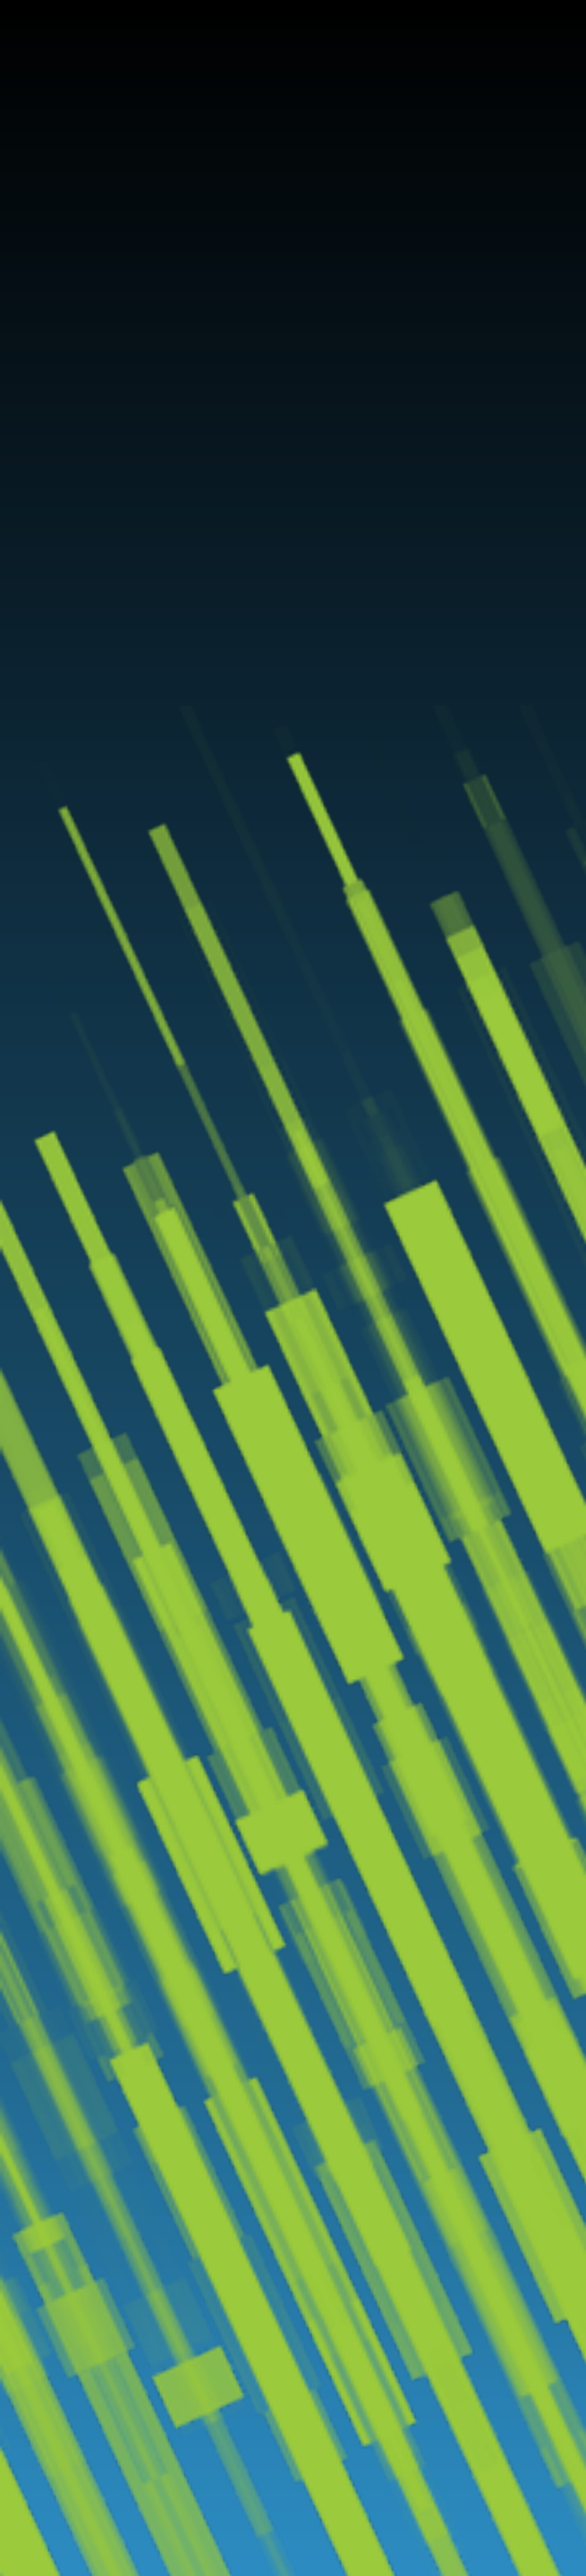
# Caching when necessary

```
def step(rdd: RDD[_]) = {  
    val wasUncached = rdd.storageLevel == StorageLevel.NONE  
    if (wasUncached) { rdd.cache() }  
    result = trainModel(rdd)  
    if (wasUncached) { rdd.unpersist() }  
  
    result  
}
```

```
var nextModel = initialModel
for (int i = 0; i < iterations; i++) {
  val current = sc.broadcast(nextModel)
  val newState = examples.aggregate(ModelState.empty()) (
    ( case (state: ModelState, example: Example) =>
      state.update(current.value.lookup(example, i), example) )
    ( case (s1: ModelState, s2: ModelState) => s1.combine(s2) )
  )
  nextModel = modelFromState(newState)
  current.unpersist
}
```

```
var nextModel = initialModel
for (int i = 0; i < iterations; i++) {
  val current = sc.broadcast(nextModel)
  val newState = examples.aggregate(ModelState.empty()) (
    ( case (state: ModelState, example: Example) =>
      state.update(current.value.lookup(example, i), example) )
    ( case (s1: ModelState, s2: ModelState) => s1.combine(s2) )
  )
  nextModel = modelFromState(newState)
  current.unpersist
}
```

```
var nextModel = initialModel
for (int i = 0; i < iterations; i++) {
  val current = sc.broadcast(nextModel)
  val newState = examples.aggregate(ModelState.empty()) (
    ( case (state: ModelState, example: Example) =>
      state.update(current.value.lookup(example, i), example) )
    ( case (s1: ModelState, s2: ModelState) => s1.combine(s2) )
  )
  nextModel = modelFromState(newState)
  current.unpersist
}
```



# Writing generic code for Spark's parallel collections



# The RDD is invariant

$T <: U \not\Rightarrow \text{RDD}[T] <: \text{RDD}[U]$

$T \leq U$   ~~$\nrightarrow$~~   $RDD[T] \leq RDD[U]$

```
trait HasUserId { val userid: Int }  
case class Transaction(override val userid: Int,  
                        timestamp: Int,  
                        amount: Double)  
    extends HasUserId {}  
  
def badKeyByUserId(r: RDD[HasUserId]) = r.map(x => (x.userid, x))
```

$T <: U$   ~~$\nrightarrow$~~   $RDD[T] <: RDD[U]$

```
trait HasUserId { val userid: Int }  
case class Transaction(override val userid: Int,  
                        timestamp: Int,  
                        amount: Double)  
    extends HasUserId {}  
  
def badKeyByUserId(r: RDD[HasUserId]) = r.map(x => (x.userid, x))
```

```
val xacts = spark.parallelize(Array(  
  Transaction(1, 1, 1.0),  
  Transaction(2, 2, 1.0)  
))
```

```
badKeyById(xacts)
```

```
<console>: error: type mismatch;
```

```
found    : org.apache.spark.rdd.RDD[Transaction]
```

```
required: org.apache.spark.rdd.RDD[HasUserId]
```

Note: Transaction <: HasUserId, but class RDD is invariant in type T.  
You may wish to define T as +T instead. (SLS 4.5)

```
badKeyById(xacts)
```

```
val xacts = spark.parallelize(Array(  
  Transaction(1, 1, 1.0),  
  Transaction(2, 2, 1.0)  
))
```

```
badKeyByUserId(xacts)
```

```
<console>: error: type mismatch;
```

```
found    : org.apache.spark.rdd.RDD[Transaction]
```

```
required: org.apache.spark.rdd.RDD[HasUserId]
```

Note: Transaction <: HasUserId, but class RDD is invariant in type T.  
You may wish to define T as +T instead. (SLS 4.5)

```
badKeyByUserId(xacts)
```



# An example: natural join

A	B	C	D	E

A	B	E	X	Y



# An example: natural join

A	B	C	D	E

A	B	E	X	Y

# An example: natural join

A	B	C	D	E	X	Y

# Ad-hoc natural join

```
df1.join(df2, df1("a") === df2("a") &&  
            df1("b") === df2("b") &&  
            df1("e") === df2("e"))
```

```

def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}

```

```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {  
    val lcols = left.columns  
    val rcols = right.columns  
    val ccols = lcols.toSet intersect rcols.toSet  
  
    if(ccols.isEmpty)  
        left.limit(0).crossJoin(right.limit(0))  
    else  
        left  
            .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))  
            .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++  
                    lcols.collect { case c if !ccols.contains(c) => left(c) } ++  
                    rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)  
}
```

```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {  
  val lcols = left.columns  
  val rcols = right.columns  
  val ccols = lcols.toSet intersect rcols.toSet  
  
  if(ccols.isEmpty)  
    left.limit(0).crossJoin(right.limit(0))  
  else  
    left  
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))  
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++  
              lcols.collect { case c if !ccols.contains(c) => left(c) } ++  
              rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)  
}
```

**introspecting over column names**





```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {  
    val lcols = left.columns  
    val rcols = right.columns  
    val ccols = lcols.toSet intersect rcols.toSet  
  
    if(ccols.isEmpty)  
        left.limit(0).crossJoin(right.limit(0))  
    else  
        left  
            .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))  
            .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++  
                    lcols.collect { case c if !ccols.contains(c) => left(c) } ++  
                    rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)  
}
```

```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } :_*)
}
```

## dynamically constructing expressions

```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {  
  val lcols = left.columns  
  val rcols = right.columns  
  val ccols = lcols.toSet intersect rcols.toSet  
  
  if(ccols.isEmpty)  
    left.limit(0).crossJoin(right.limit(0))  
  else  
    left  
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))  
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++  
              lcols.collect { case c if !ccols.contains(c) => left(c) } ++  
              rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)  
}
```

## dynamically constructing expressions

```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}
```

## dynamically constructing expressions

```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}
```

**[left.a === right.a, left.b === right.b, ...]**



```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}
```

**left.a === right.a && left.b === right.b && ...**



```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}
```

**left.a === right.a && left.b === right.b && ...**





```

def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}

```



```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}
```

**dynamically constructing column lists**

```
def natjoin(left: DataFrame, right: DataFrame): DataFrame = {
  val lcols = left.columns
  val rcols = right.columns
  val ccols = lcols.toSet intersect rcols.toSet

  if(ccols.isEmpty)
    left.limit(0).crossJoin(right.limit(0))
  else
    left
      .join(right, ccols.map {col => left(col) === right(col) }.reduce(_ && _))
      .select(lcols.collect { case c if ccols.contains(c) => left(c) } ++
        lcols.collect { case c if !ccols.contains(c) => left(c) } ++
        rcols.collect { case c if !ccols.contains(c) => right(c) } : _*)
}
```

**dynamically constructing column lists**

# User-defined functions

```
{"a": 1, "b": "wilma", ..., "x": "club"}  
{"a": 2, "b": "betty", ..., "x": "diamond"}  
{"a": 3, "b": "fred", ..., "x": "heart"}  
{"a": 4, "b": "barney", ..., "x": "spade"}
```

# User-defined functions

```
{"a": 1, "b": "wilma", ..., "x": "club"}  
{"a": 2, "b": "betty", ..., "x": "diamond"}  
{"a": 3, "b": "fred", ..., "x": "heart"}  
{"a": 4, "b": "barney", ..., "x": "spade"}
```



wilma	club
betty	diamond
fred	heart
barney	spade

```
import json
from pyspark.sql.types import *
from pyspark.sql.functions import udf

def selectively_structure(fields):
    resultType = StructType([StructField(f, StringType(), nullable=True)
                              for f in fields])
    def impl(js):
        try:
            d = json.loads(js)
            return [str(d.get(f)) for f in fields]
        except:
            return [None] * len(fields)
    return udf(impl, resultType)
```

```
import json
from pyspark.sql.types import *
from pyspark.sql.functions import udf

def selectively_structure(fields):
    resultType = StructType([StructField(f, StringType(), nullable=True)
                              for f in fields])
    def impl(js):
        try:
            d = json.loads(js)
            return [str(d.get(f)) for f in fields]
        except:
            return [None] * len(fields)
    return udf(impl, resultType)
```

```
import json
from pyspark.sql.types import *
from pyspark.sql.functions import udf

def selectively_structure(fields):
    resultType = StructType([StructField(f, StringType(), nullable=True)
                              for f in fields])
    def impl(js):
        try:
            d = json.loads(js)
            return [str(d.get(f)) for f in fields]
        except:
            return [None] * len(fields)
    return udf(impl, resultType)
```

```
import json
from pyspark.sql.types import *
from pyspark.sql.functions import udf

def selectively_structure(fields):
    resultType = StructType([StructField(f, StringType(), nullable=True)
                              for f in fields])
    def impl(js):
        try:
            d = json.loads(js)
            return [str(d.get(f)) for f in fields]
        except:
            return [None] * len(fields)
    return udf(impl, resultType)
```



```
import json
from pyspark.sql.types import *
from pyspark.sql.functions import udf

def selectively_structure(fields):
    resultType = StructType([StructField(f, StringType(), nullable=True)
                              for f in fields])
    def impl(js):
        try:
            d = json.loads(js)
            return [str(d.get(f)) for f in fields]
        except:
            return [None] * len(fields)
    return udf(impl, resultType)
```

```
import json
from pyspark.sql.types import *
from pyspark.sql.functions import udf


def selectively_structure(fields):
    resultType = StructType([StructField(f, StringType(), nullable=True)
        for f in fields])
    def impl(js):
        try:
            d = json.loads(js)
            return [str(d.get(f)) for f in fields]
        except:
            return [None] * len(fields)
    return udf(impl, resultType)

extract_bx = selectively_structure(["b", "x"])

structured_df = df.withColumn("result", extract_bx("json"))
```


# Spark's ML pipelines

# Spark's ML pipelines




`model.transform(df)`

# Spark's ML pipelines




estimator.**fit**(df)



model.**transform**(df)

# Spark's ML pipelines



# Building Machine Learning Algorithms on Apache Spark: Scaling Out and Up

---

There are lots of reasons why you might want to implement your own machine learning algorithms on Spark: you might want to experiment with a new idea, try and reproduce results from a recent research paper, or simply to use an existing technique that isn't implemented in MLlib.

In this talk, we'll walk through the process of developing a new machine learning algorithm for Spark. We'll start with the basics, by considering how we'd design a scale-out parallel implementation of our unsupervised learning technique. The bulk of the talk will focus on the details you need to know to turn an algorithm design into an efficient parallel implementation on Spark.

We'll start by reviewing a simple RDD-based implementation, show some improvements, point out some pitfalls to avoid, and iteratively extend our implementation to support contemporary Spark features like ML Pipelines and structured query processing. We'll conclude by briefly examining some useful techniques to complement scale-out performance by scaling our code up, taking advantage of specialized hardware to accelerate single-worker performance.

You'll leave this talk with everything you need to build a new machine learning technique that runs on Spark.

Session hashtag: #DS4SAIS





















# User-defined aggregates: the fundamentals









# Three components








# Three components









# Three components



# Three components



# Three components





# User-defined aggregates: the implementation

```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
  (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
extends UserDefinedAggregateFunction {

  def deterministic: Boolean = false

  def inputSchema: StructType =
    StructType(StructField("x", dataType.tpe) :: Nil)

  def bufferSchema: StructType =
    StructType(StructField("tdigest", TDigestUDT) :: Nil)

  def dataType: DataType = TDigestUDT
```

```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
  (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
  extends UserDefinedAggregateFunction {

  def deterministic: Boolean = false

  def inputSchema: StructType =
    StructType(StructField("x", dataType.tpe) :: Nil)

  def bufferSchema: StructType =
    StructType(StructField("tdigest", TDigestUDT) :: Nil)

  def dataType: DataType = TDigestUDT
```



```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
    (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
    extends UserDefinedAggregateFunction {

    def deterministic: Boolean = false

    def inputSchema: StructType =
        StructType(StructField("x", dataType.tpe) :: Nil)

    def bufferSchema: StructType =
        StructType(StructField("tdigest", TDigestUDT) :: Nil)

    def dataType: DataType = TDigestUDT
```

```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
  (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
  extends UserDefinedAggregateFunction {

  def deterministic: Boolean = false

  def inputSchema: StructType =
    StructType(StructField("x", dataType.tpe) :: Nil)

  def bufferSchema: StructType =
    StructType(StructField("tdigest", TDigestUDT) :: Nil)

  def dataType: DataType = TDigestUDT
```

```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
  (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
  extends UserDefinedAggregateFunction {

  def deterministic: Boolean = false

  def inputSchema: StructType =
    StructType(StructField("x", dataType.tpe) :: Nil)

  def bufferSchema: StructType =
    StructType(StructField("tdigest", TDigestUDT) :: Nil)

  def dataType: DataType = TDigestUDT
```

```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
  (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
  extends UserDefinedAggregateFunction {

  def deterministic: Boolean = false

  def inputSchema: StructType =
    StructType(StructField("x", dataType.tpe) :: Nil)

  def bufferSchema: StructType =
    StructType(StructField("tdigest", TDigestUDT) :: Nil)

  def dataType: DataType = TDigestUDT
```

```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
    (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
    extends UserDefinedAggregateFunction {

    def deterministic: Boolean = false

    def inputSchema: StructType =
        StructType(StructField("x", dataType.tpe) :: Nil)

    def bufferSchema: StructType =
        StructType(StructField("tdigest", TDigestUDT) :: Nil)

    def dataType: DataType = TDigestUDT
```

```
case class TDigestUDAF[N](deltaV: Double, maxDiscreteV: Int)
  (implicit num: Numeric[N], dataType: TDigestUDAFDataType[N])
  extends UserDefinedAggregateFunction {


  def deterministic: Boolean = false

  def inputSchema: StructType =
    StructType(StructField("x", dataType.tpe) :: Nil)

  def bufferSchema: StructType =
    StructType(StructField("tdigest", TDigestUDT) :: Nil)

  def dataType: DataType = TDigestUDT
```

# Four main functions: initialize



`initialize`

# Four main functions: initialize



**initialize**



```
def initialize(buf: MutableAggregationBuffer): Unit = {  
    buf(0) = TDigestSQL(TDigest.empty(deltaV, maxDiscreteV))  
}  
  
def evaluate(buf: Row): Any = buf.getAs[TDigestSQL](0)
```

```
def initialize(buf: MutableAggregationBuffer): Unit = {  
    buf(0) = TDigestSQL(TDigest.empty(deltaV, maxDiscreteV))  
}  
  
def evaluate(buf: Row): Any = buf.getAs[TDigestSQL](0)
```

# Four main functions: evaluate



evaluate

# Four main functions: evaluate



evaluate

```
def initialize(buf: MutableAggregationBuffer): Unit = {  
    buf(0) = TDigestSQL(TDigest.empty(deltaV, maxDiscreteV))  
}  
  
def evaluate(buf: Row): Any = buf.getAs[TDigestSQL](0)
```

```
def initialize(buf: MutableAggregationBuffer): Unit = {  
    buf(0) = TDigestSQL(TDigest.empty(deltaV, maxDiscreteV))  
}  
  
def evaluate(buf: Row): Any = buf.getAs[TDigestSQL](0)
```

# Four main functions: update



update

# Four main functions: update



update



```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
  if (!input.isNullAt(0)) {  
    buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
      num.toDouble(input.getAs[N](0)))  
  }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
  buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
    buf2.getAs[TDigestSQL](0).tdigest)  
}
```

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```

# Four main functions: merge



**merge**



# Four main functions: merge



1 + 2

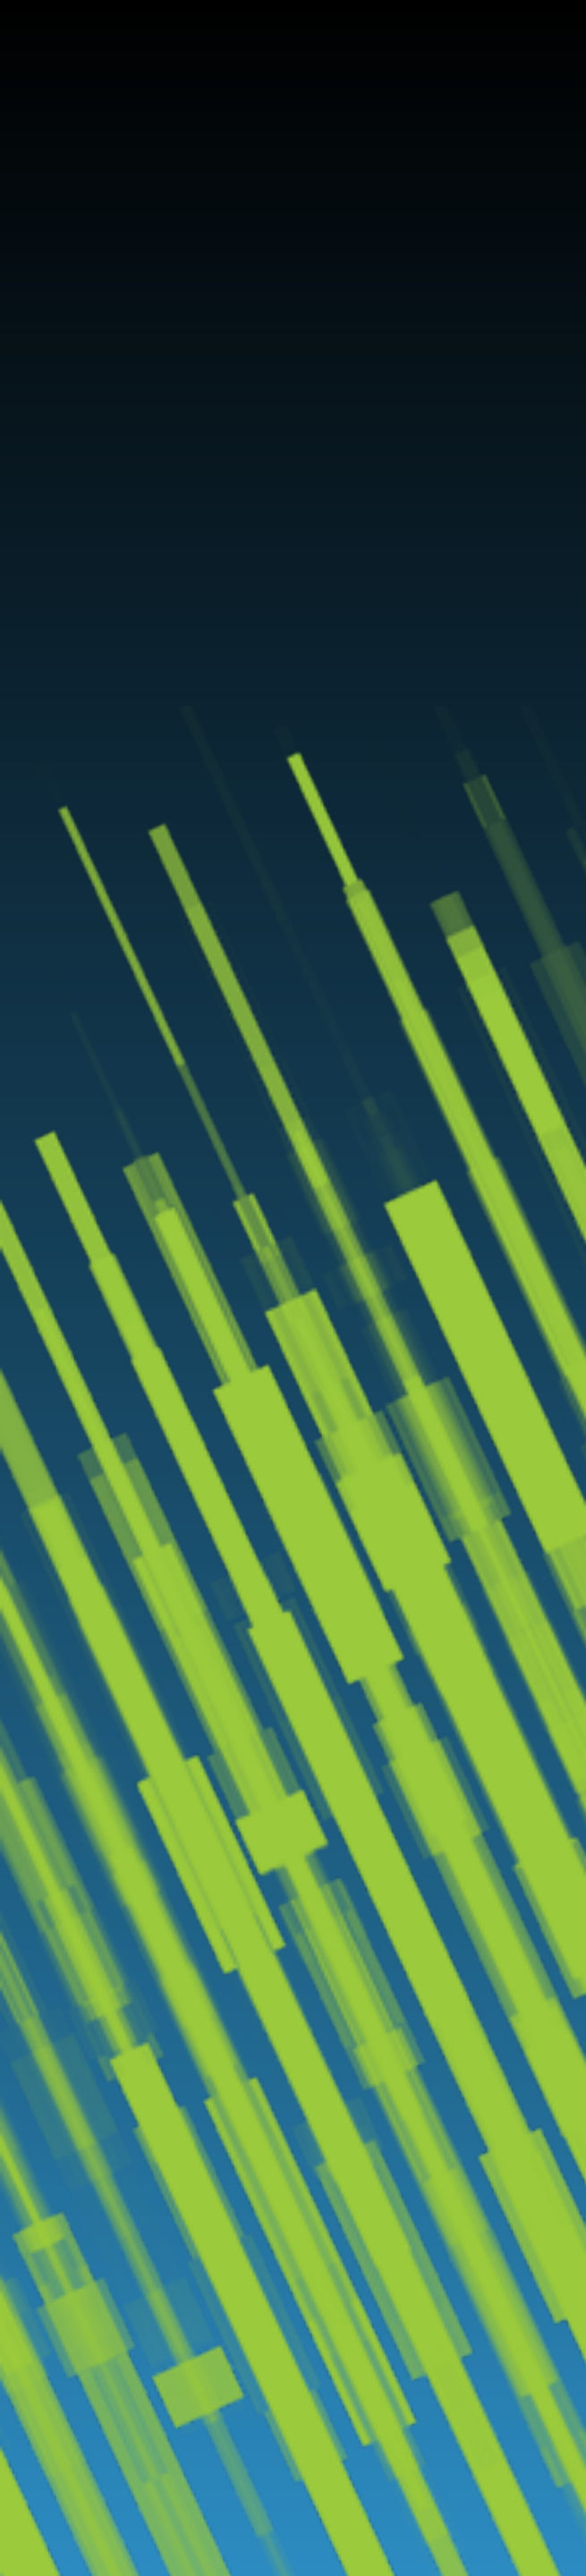
merge

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```

```
def update(buf: MutableAggregationBuffer, input: Row): Unit = {  
    if (!input.isNullAt(0)) {  
        buf(0) = TDigestSQL(buf.getAs[TDigestSQL](0).tdigest +  
            num.toDouble(input.getAs[N](0)))  
    }  
}
```

```
def merge(buf1: MutableAggregationBuffer, buf2: Row): Unit = {  
    buf1(0) = TDigestSQL(buf1.getAs[TDigestSQL](0).tdigest ++  
        buf2.getAs[TDigestSQL](0).tdigest)  
}
```



**User-defined aggregates:  
User-defined types**

# User-defined types

```
package org.apache.spark.isarnproject.sketches.udt

@SQLUserDefinedType(udt = classOf[TDigestUDT])
case class TDigestSQL(tdigest: TDigest)

class TDigestUDT extends UserDefinedType[TDigestSQL] {
  def userClass: Class[TDigestSQL] = classOf[TDigestSQL]

  // ...
```

# User-defined types

```
package org.apache.spark.isarnproject.sketches.udt

@SQLUserDefinedType(udt = classOf[TDigestUDT])
case class TDigestSQL(tdigest: TDigest)

class TDigestUDT extends UserDefinedType[TDigestSQL] {
  def userClass: Class[TDigestSQL] = classOf[TDigestSQL]

  // ....
```

# User-defined types

```
package org.apache.spark.isarnproject.sketches.udt

@SQLUserDefinedType(udt = classOf[TDigestUDT])
case class TDigestSQL(tdigest: TDigest)

class TDigestUDT extends UserDefinedType[TDigestSQL] {
  def userClass: Class[TDigestSQL] = classOf[TDigestSQL]

  // ....
```

# Implementing custom types

```
class TDigestUDT extends UserDefinedType[TDigestSQL] {  
  def userClass: Class[TDigestSQL] = classOf[TDigestSQL]  
  
  override def pyUDT: String =  
    "isarnproject.sketches.udt.tdigest.TDigestUDT"  
  
  override def typeName: String = "tdigest"  
  
  def sqlType: DataType = StructType(  
    StructField("delta", DoubleType, false) ::  
    /* ... */  
    StructField("clustM", ArrayType(DoubleType, false), false) ::  
    Nil)
```



```
class TDigestUDT extends UserDefinedType[TDigestSQL] {  
  def userClass: Class[TDigestSQL] = classOf[TDigestSQL]  
  
  override def pyUDT: String =  
    "isarnproject.sketches.udt.tdigest.TDigestUDT"  
  
  override def typeName: String = "tdigest"  
  
  def sqlType: DataType = StructType(  
    StructField("delta", DoubleType, false) ::  
    /* ... */  
    StructField("clustM", ArrayType(DoubleType, false), false) ::  
    Nil)
```

```
class TDigestUDT extends UserDefinedType[TDigestSQL] {  
  def userClass: Class[TDigestSQL] = classOf[TDigestSQL]  
  
  override def pyUDT: String =  
    "isarnproject.sketches.udt.tdigest.TDigestUDT"  
  
  override def typeName: String = "tdigest"  
  
  def sqlType: DataType = StructType(  
    StructField("delta", DoubleType, false) ::  
    /* ... */  
    StructField("clustM", ArrayType(DoubleType, false), false) ::  
    Nil)
```

```
class TDigestUDT extends UserDefinedType[TDigestSQL] {  
  def userClass: Class[TDigestSQL] = classOf[TDigestSQL]  
  
  override def pyUDT: String =  
    "isarnproject.sketches.udt.tdigest.TDigestUDT"  
  
  override def typeName: String = "tdigest"  
  
  def sqlType: DataType = StructType(  
    StructField("delta", DoubleType, false) ::  
    /* ... */  
    StructField("clustM", ArrayType(DoubleType, false), false) ::  
    Nil)
```

```
def serialize(tdsq1: TDigestSQL): Any = serializeTD(tdsq1.tdigest)

private[sketches] def serializeTD(td: TDigest): InternalRow = {
  val TDigest(delta, maxDiscrete, nclusters, clusters) = td
  val row = new GenericInternalRow(5)
  row.setDouble(0, delta)
  row.setInt(1, maxDiscrete)
  row.setInt(2, nclusters)
  val clustX = clusters.keys.toArray
  val clustM = clusters.values.toArray
  row.update(3, UnsafeArrayData.fromPrimitiveArray(clustX))
  row.update(4, UnsafeArrayData.fromPrimitiveArray(clustM))
  row
}
```



```
def serialize(tdsq1: TDigestSQL): Any = serializeTD(tdsq1.tdigest)

private[sketches] def serializeTD(td: TDigest): InternalRow = {
  val TDigest(delta, maxDiscrete, nclusters, clusters) = td
  val row = new GenericInternalRow(5)
  row.setDouble(0, delta)
  row.setInt(1, maxDiscrete)
  row.setInt(2, nclusters)
  val clustX = clusters.keys.toArray
  val clustM = clusters.values.toArray
  row.update(3, UnsafeArrayData.fromPrimitiveArray(clustX))
  row.update(4, UnsafeArrayData.fromPrimitiveArray(clustM))
  row
}
```

```
def serialize(tdsq1: TDigestSQL): Any = serializeTD(tdsq1.tdigest)

private[sketches] def serializeTD(td: TDigest): InternalRow = {
  val TDigest(delta, maxDiscrete, nclusters, clusters) = td
  val row = new GenericInternalRow(5)
  row.setDouble(0, delta)
  row.setInt(1, maxDiscrete)
  row.setInt(2, nclusters)
  val clustX = clusters.keys.toArray
  val clustM = clusters.values.toArray
  row.update(3, UnsafeArrayData.fromPrimitiveArray(clustX))
  row.update(4, UnsafeArrayData.fromPrimitiveArray(clustM))
  row
}
```

```
def serialize(tdsq1: TDigestSQL): Any = serializeTD(tdsq1.tdigest)

private[sketches] def serializeTD(td: TDigest): InternalRow = {
  val TDigest(delta, maxDiscrete, nclusters, clusters) = td
  val row = new GenericInternalRow(5)
  row.setDouble(0, delta)
  row.setInt(1, maxDiscrete)
  row.setInt(2, nclusters)
  val clustX = clusters.keys.toArray
  val clustM = clusters.values.toArray
  row.update(3, UnsafeArrayData.fromPrimitiveArray(clustX))
  row.update(4, UnsafeArrayData.fromPrimitiveArray(clustM))
  row
}
```

```
def deserialize(td: Any): TDigestSQL = TDigestSQL(deserializeTD(td))
```

```
private[sketches] def deserializeTD(datum: Any): TDigest =  
  datum match { case row: InternalRow =>  
    val delta = row.getDouble(0)  
    val maxDiscrete = row.getInt(1)  
    val nclusters = row.getInt(2)  
    val clustX = row.getArray(3).toDoubleArray()  
    val clustM = row.getArray(4).toDoubleArray()  
    val clusters = clustX.zip(clustM)  
      .foldLeft(TDigestMap.empty) { case (td, e) => td + e }  
    TDigest(delta, maxDiscrete, nclusters, clusters)  
  }
```



```
def deserialize(td: Any): TDigestSQL = TDigestSQL(deserializeTD(td))


private[sketches] def deserializeTD(datum: Any): TDigest =
  datum match { case row: InternalRow =>
    val delta = row.getDouble(0)
    val maxDiscrete = row.getInt(1)
    val nclusters = row.getInt(2)
    val clustX = row.getDoubleArray(3)
    val clustM = row.getDoubleArray(4)
    val clusters = clustX.zip(clustM)
      .foldLeft(TDigestMap.empty) { case (td, e) => td + e }
    TDigest(delta, maxDiscrete, nclusters, clusters)
  }
```


```
def deserialize(td: Any): TDigestSQL = TDigestSQL(deserializeTD(td))



private[sketches] def deserializeTD(datum: Any): TDigest =
  datum match { case row: InternalRow =>
    val delta = row.getDouble(0)
    val maxDiscrete = row.getInt(1)
    val nclusters = row.getInt(2)
    val clustX = row.getDoubleArray(3)
    val clustM = row.getDoubleArray(4)
    val clusters = clustX.zip(clustM)
      .foldLeft(TDigestMap.empty) { case (td, e) => td + e }
    TDigest(delta, maxDiscrete, nclusters, clusters)
  }
```







# Extending PySpark with your Scala library













```
# class to access the active Spark context for Python  
from pyspark.context import SparkContext
```

```
# gateway to the JVM from py4j  
sparkJVM = SparkContext._active_spark_context._jvm
```

```
# use the gateway to access JVM objects and classes  
thisThing = sparkJVM.com.path.to.this.thing
```







```
# class to access the active Spark context for Python  
from pyspark.context import SparkContext
```

```
# gateway to the JVM from py4j  
sparkJVM = SparkContext._active_spark_context._jvm
```

```
# use the gateway to access JVM objects and classes  
thisThing = sparkJVM.com.path.to.this.thing
```



```
# class to access the active Spark context for Python  
from pyspark.context import SparkContext
```

```
# gateway to the JVM from py4j  
sparkJVM = SparkContext._active_spark_context._jvm
```

```
# use the gateway to access JVM objects and classes  
thisThing = sparkJVM.com.path.to.this.thing
```

# A Python-friendly wrapper

```
package org.isarnproject.sketches.udaf
```

```
object pythonBindings {
```

```
  def tdigestDoubleUDAF(delta: Double, maxDiscrete: Int) =
```

```
    TDigestUDAF[Double](delta, maxDiscrete)
```

```
}
```

```
package org.isarnproject.sketches.udaf

object pythonBindings {
  def tdigestDoubleUDAF(delta: Double, maxDiscrete: Int) =
    TDigestUDAF[Double](delta, maxDiscrete)
}
```

```
package org.isarnproject.sketches.udaf

object pythonBindings {
  def tdigestDoubleUDAF(delta: Double, maxDiscrete: Int) =
    TDigestUDAF[Double](delta, maxDiscrete)
}
```

```
from pyspark.sql.column import Column, _to_java_column, _to_seq
from pyspark.context import SparkContext

# one of these for each type parameter Double, Int, Long, etc
def tdigestDoubleUDAF(col, delta=0.5, maxDiscrete=0):
    sc = SparkContext._active_spark_context
    pb = sc._jvm.org.isarnproject.sketches.udaf.pythonBindings
    tdapply = pb.tdigestDoubleUDAF(delta, maxDiscrete).apply
    return Column(tdapply(_to_seq(sc, [col], _to_java_column)))
```

```
from pyspark.sql.column import Column, _to_java_column, _to_seq
from pyspark.context import SparkContext

# one of these for each type parameter Double, Int, Long, etc
def tdigestDoubleUDAF(col, delta=0.5, maxDiscrete=0):
    sc = SparkContext._active_spark_context
    pb = sc._jvm.org.isarnproject.sketches.udaf.pythonBindings
    tdapply = pb.tdigestDoubleUDAF(delta, maxDiscrete).apply
    return Column(tdapply(_to_seq(sc, [col], _to_java_column)))
```

```
from pyspark.sql.column import Column, _to_java_column, _to_seq
from pyspark.context import SparkContext

# one of these for each type parameter Double, Int, Long, etc
def tdigestDoubleUDAF(col, delta=0.5, maxDiscrete=0):
    sc = SparkContext._active_spark_context
    pb = sc._jvm.org.isarnproject.sketches.udaf.pythonBindings
    tdapply = pb.tdigestDoubleUDAF(delta, maxDiscrete).apply
    return Column(tdapply(_to_seq(sc, [col], _to_java_column)))
```



```
class TDigestUDT(UserDefinedType):
    @classmethod
    def sqlType(cls):
        return StructType([
            StructField("delta", DoubleType(), False),
            StructField("maxDiscrete", IntegerType(), False),
            StructField("nclusters", IntegerType(), False),
            StructField("clustX", ArrayType(DoubleType(), False), False),
            StructField("clustM", ArrayType(DoubleType(), False), False)]

# ...
```

```
class TDigestUDT(UserDefinedType):
    @classmethod
    def sqlType(cls):
        return StructType([
            StructField("delta", DoubleType(), False),
            StructField("maxDiscrete", IntegerType(), False),
            StructField("nclusters", IntegerType(), False),
            StructField("clustX", ArrayType(DoubleType(), False), False),
            StructField("clustM", ArrayType(DoubleType(), False), False)])

# ...
```

```
class TDigestUDT(UserDefinedType):  
    # ...  
    @classmethod  
    def module(cls):  
        return "isarnproject.sketches.udt.tdigest"  
  
    @classmethod  
    def scalaUDT(cls):  
        return "org.apache.spark.isarnproject.sketches.udt.TDigestUDT"  
  
    def simpleString(self):  
        return "tdigest"
```

```
class TDigestUDT(UserDefinedType):  
    # ...  
    @classmethod  
    def module(cls):  
        return "isarnproject.sketches.udt.tdigest"  
  
    @classmethod  
    def scalaUDT(cls):  
        return "org.apache.spark.isarnproject.sketches.udt.TDigestUDT"  
  
    def simpleString(self):  
        return "tdigest"
```

```
class TDigestUDT(UserDefinedType):  
    # ...  
    @classmethod  
    def module(cls):  
        return "isarnproject.sketches.udt.tdigest"  
  
    @classmethod  
    def scalaUDT(cls):  
        return "org.apache.spark.isarnproject.sketches.udt.TDigestUDT"  
  
    def simpleString(self):  
        return "tdigest"
```

```
class TDigestUDT(UserDefinedType):  
    # ...  
    def serialize(self, obj):  
        return (obj.delta, obj.maxDiscrete, obj.nclusters, \  
                [float(v) for v in obj.clustX], \  
                [float(v) for v in obj.clustM])  
  
    def deserialize(self, datum):  
        return TDigest(datum[0], datum[1], datum[2], datum[3], datum[4])
```

```
class TDigestUDT(UserDefinedType):
    # ...

    def serialize(self, obj):
        return (obj.delta, obj.maxDiscrete, obj.nclusters, \
                [float(v) for v in obj.clustX], \
                [float(v) for v in obj.clustM])

    def deserialize(self, datum):
        return TDigest(datum[0], datum[1], datum[2], datum[3], datum[4])
```

```
class TDigestUDT extends UserDefinedType[TDigestSQL] {  
  // ...  
  override def pyUDT: String =  
    “isarnproject.sketches.udt.tdigest.TDigestUDT”  
}
```



# Python code in JAR files

```
mappings in (Compile, packageBin) += Seq(  
  (baseDirectory.value / "python" / "isarnproject" / "__init__.pyc") ->  
    "isarnproject/__init__.pyc",  
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "__init__.pyc") ->  
    "isarnproject/sketches/__init__.pyc",  
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "__init__.pyc") ->  
    "isarnproject/sketches/udaf/__init__.pyc",  
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "tdigest.pyc") ->  
    "isarnproject/sketches/udaf/tdigest.pyc",  
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "__init__.pyc") ->  
    "isarnproject/sketches/udt/__init__.pyc",  
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "tdigest.pyc") ->  
    "isarnproject/sketches/udt/tdigest.pyc"  
)
```

```
mappings in (Compile, packageBin) += Seq(
  (baseDirectory.value / "python" / "isarnproject" / "__init__.pyc") ->
    "isarnproject/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "__init__.pyc") ->
    "isarnproject/sketches/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "__init__.pyc") ->
    "isarnproject/sketches/udaf/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "tdigest.pyc") ->
    "isarnproject/sketches/udaf/tdigest.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "__init__.pyc") ->
    "isarnproject/sketches/udt/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "tdigest.pyc") ->
    "isarnproject/sketches/udt/tdigest.pyc"
)
```

```
mappings in (Compile, packageBin) += Seq(
  (baseDirectory.value / "python" / "isarnproject" / "__init__.pyc") ->
    "isarnproject/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "__init__.pyc") ->
    "isarnproject/sketches/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "__init__.pyc") ->
    "isarnproject/sketches/udaf/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "tdigest.pyc") ->
    "isarnproject/sketches/udaf/tdigest.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "__init__.pyc") ->
    "isarnproject/sketches/udt/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "tdigest.pyc") ->
    "isarnproject/sketches/udt/tdigest.pyc"
)
```



```
mappings in (Compile, packageBin) += Seq(
  (baseDirectory.value / "python" / "isarnproject" / "__init__.pyc") ->
    "isarnproject/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "__init__.pyc") ->
    "isarnproject/sketches/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "__init__.pyc") ->
    "isarnproject/sketches/udaf/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udaf" / "tdigest.pyc") ->
    "isarnproject/sketches/udaf/tdigest.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "__init__.pyc") ->
    "isarnproject/sketches/udt/__init__.pyc",
  (baseDirectory.value / "python" / "isarnproject" / "sketches" / "udt" / "tdigest.pyc") ->
    "isarnproject/sketches/udt/tdigest.pyc"
)
```



# Cross-building for Python

```
lazy val compilePython = taskKey[Unit]("Compile python files")
```

```
compilePython := {  
  val s: TaskStreams = streams.value  
  s.log.info("compiling python...")  
  val stat = (Seq(pythonCMD, "-m", "compileall", "python/") !)  
  if (stat != 0) {  
    throw new IllegalStateException("python compile failed")  
  }  
}
```

```
(packageBin in Compile) <=<  
  (packageBin in Compile).dependsOn(compilePython)
```

```
lazy val compilePython = taskKey[Unit]("Compile python files")
```

```
compilePython := {  
  val s: TaskStreams = streams.value  
  s.log.info("compiling python...")  
  val stat = (Seq(pythonCMD, "-m", "compileall", "python/") !)  
  if (stat != 0) {  
    throw new IllegalStateException("python compile failed")  
  }  
}
```

```
(packageBin in Compile) <=<  
  (packageBin in Compile).dependsOn(compilePython)
```

```
lazy val compilePython = taskKey[Unit]("Compile python files")

compilePython := {
  val s: TaskStreams = streams.value
  s.log.info("compiling python...")
  val stat = (Seq(pythonCMD, "-m", "compileall", "python/") !)
  if (stat != 0) {
    throw new IllegalStateException("python compile failed")
  }
}

(packageBin in Compile) <=<=
  (packageBin in Compile).dependsOn(compilePython)
```

# Using versioned JAR files

```
$ pyspark --packages \  
    'org.isarnproject:isarn-sketches-spark_2.11:0.3.0-sp2.2-py2.7'
```



# Using versioned JAR files

```
$ pyspark --packages \
    'org.isarnproject:isarn-sketches-spark_2.11:0.3.0-sp2.2-py2.7'
```

# Using versioned JAR files

```
$ pyspark --packages \
    'org.isarnproject:isarn-sketches-spark_2.11:0.3.0-sp2.2-py2.7'
```



**Show your work:  
publishing results**

# Maven Central

# Bintray

not really

**easy to set up for library developers**

trivial

trivial

**easy to set up for library users**

mostly

yes, via sbt

**easy to publish**

yes, via sbt + plugins







yes

**easy to resolve artifacts**


mostly







# Conclusions and takeaways








`estimator.fit(df)`




`model.transform(df)`











estimator.**fit**(df)




model.**transform**(df)













estimator.**fit**(df)




model.**transform**(df)













estimator.**fit**(df)



model.**transform**(df)





# KEEP IN TOUCH

<https://radanalytics.io>

[willb@redhat.com](mailto:willb@redhat.com) • @willb

[eje@redhat.com](mailto:eje@redhat.com) • @manyangled