

Automated Debugging of Big Data Analytics in Apache Spark Using BigSift

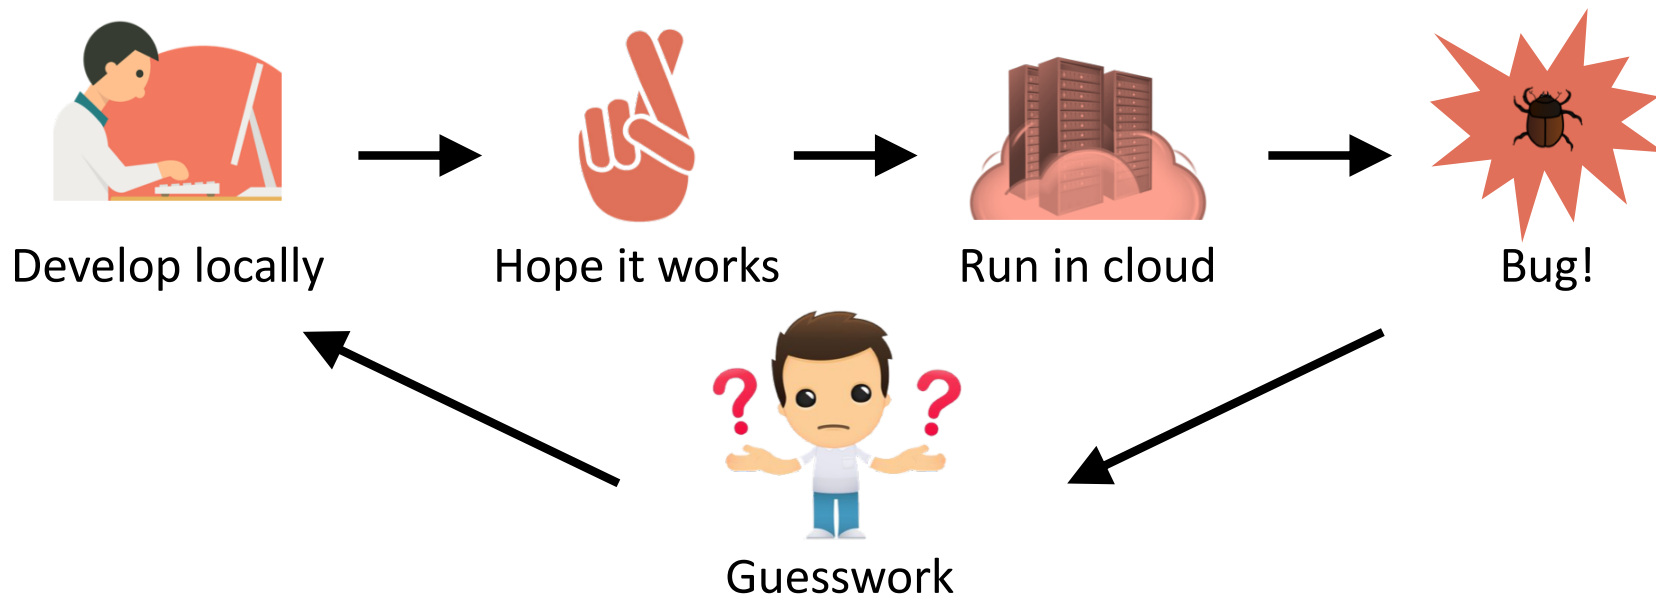
Muhammad Ali Gulzar Miryung Kim
University of California, Los Angeles

The UCLA logo, consisting of the letters 'UCLA' in white on a blue square background.

UCLA

#Res3SAIS

Big Data Debugging in the Dark



Google
Map Reduce



Spark



BigDebug: Debugging Primitives for Interactive Big Data Processing

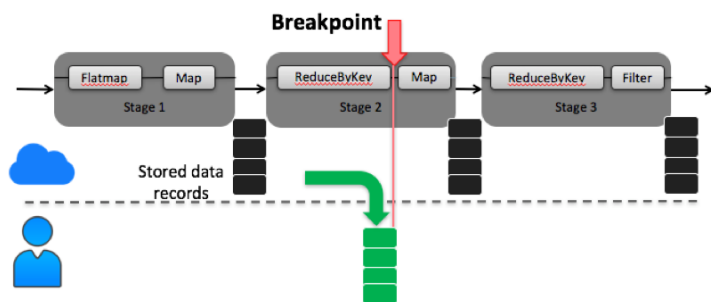
ICSE 2016

Muhammad Ali Gulzar, Matteo Interlandi, Seunghyun Yoo, Sai Deep Tetali Tyson Condie, Todd Millstein, Miryung Kim

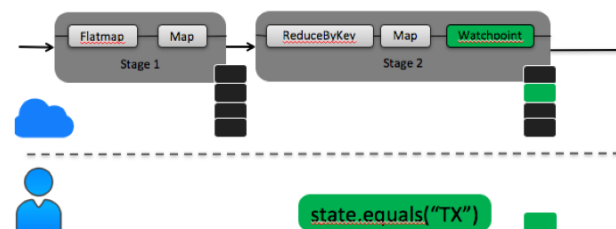
Presented at [Spark Summit 2017](#)

Interactive Debugging with BigDebug

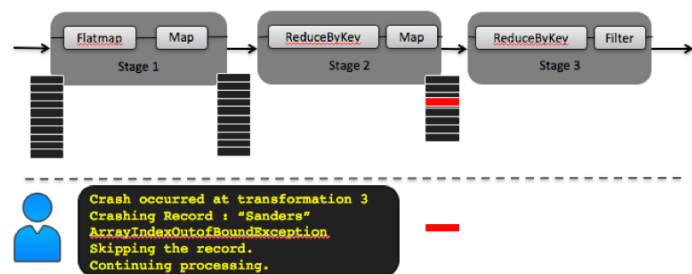
1. Simulated Breakpoint



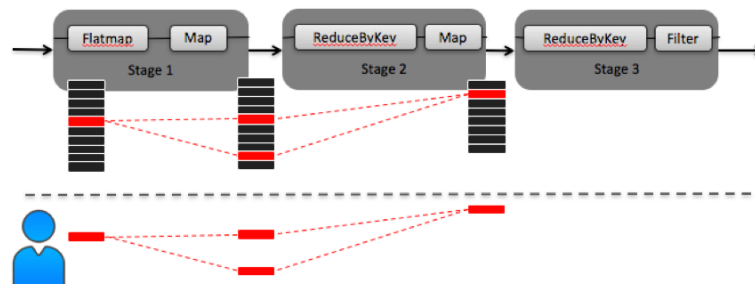
2. On Demand Guarded Watchpoint



3. Crash Culprit Identification



4. Backward and Forward Tracing



Titian: Data Provenance support in Spark

VLDB 2015

Matteo Interlandi, Kshitij Shah, Sai Deep Tetali, Muhammad Ali Gulzar, Seunghyun Yoo, Miryung Kim, Todd Millstein, Tyson Condie

Data Provenance – in SQL

```
SELECT time, AVG(temp)
FROM sensors
GROUP BY time
```

Why ID-2 and
ID-3 have those
high values?

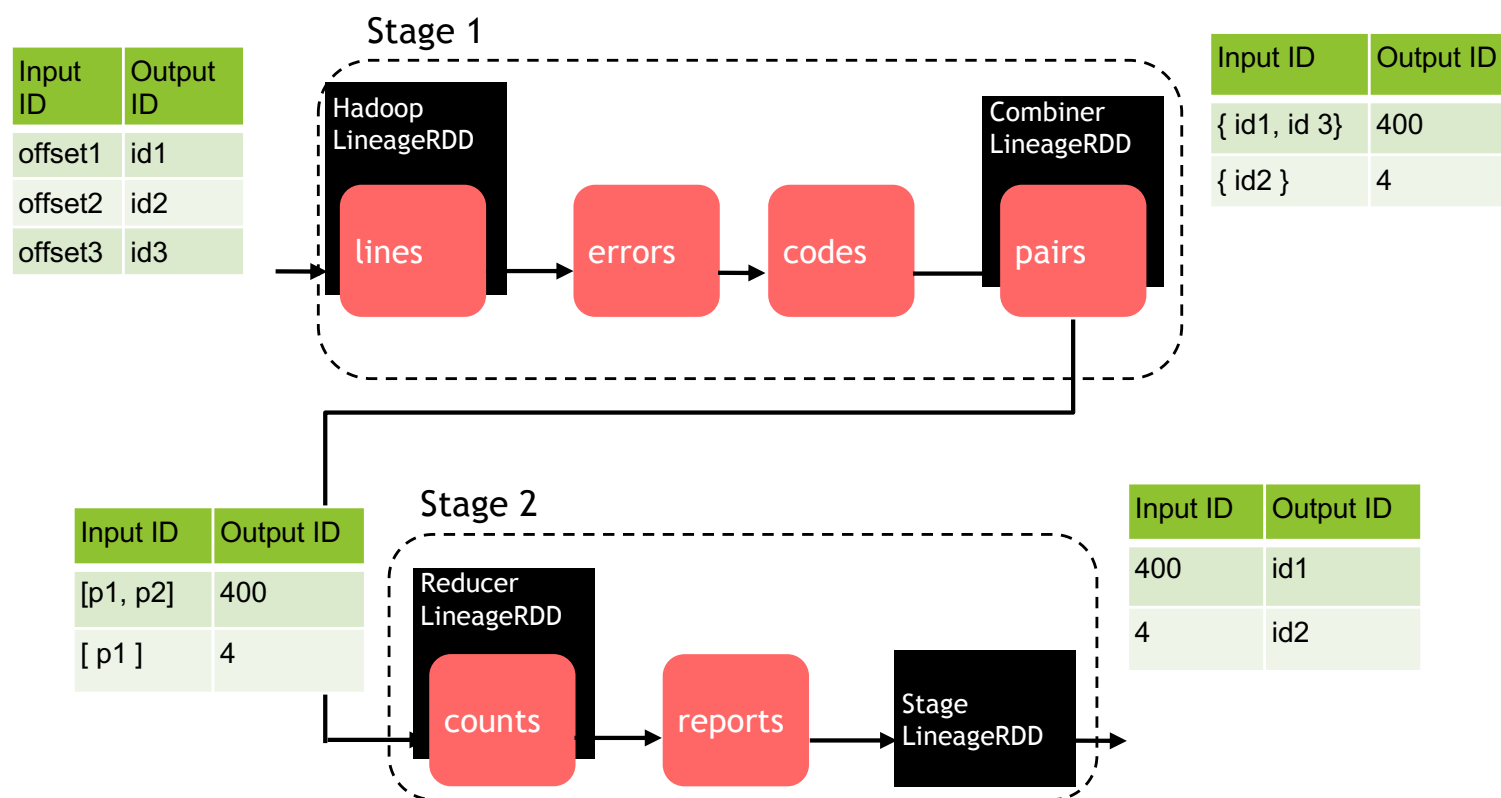
Result-ID	Time	AVG(temp)
ID-1	11AM	34.6
ID-2	12PM	56.6
ID-3	1PM	50



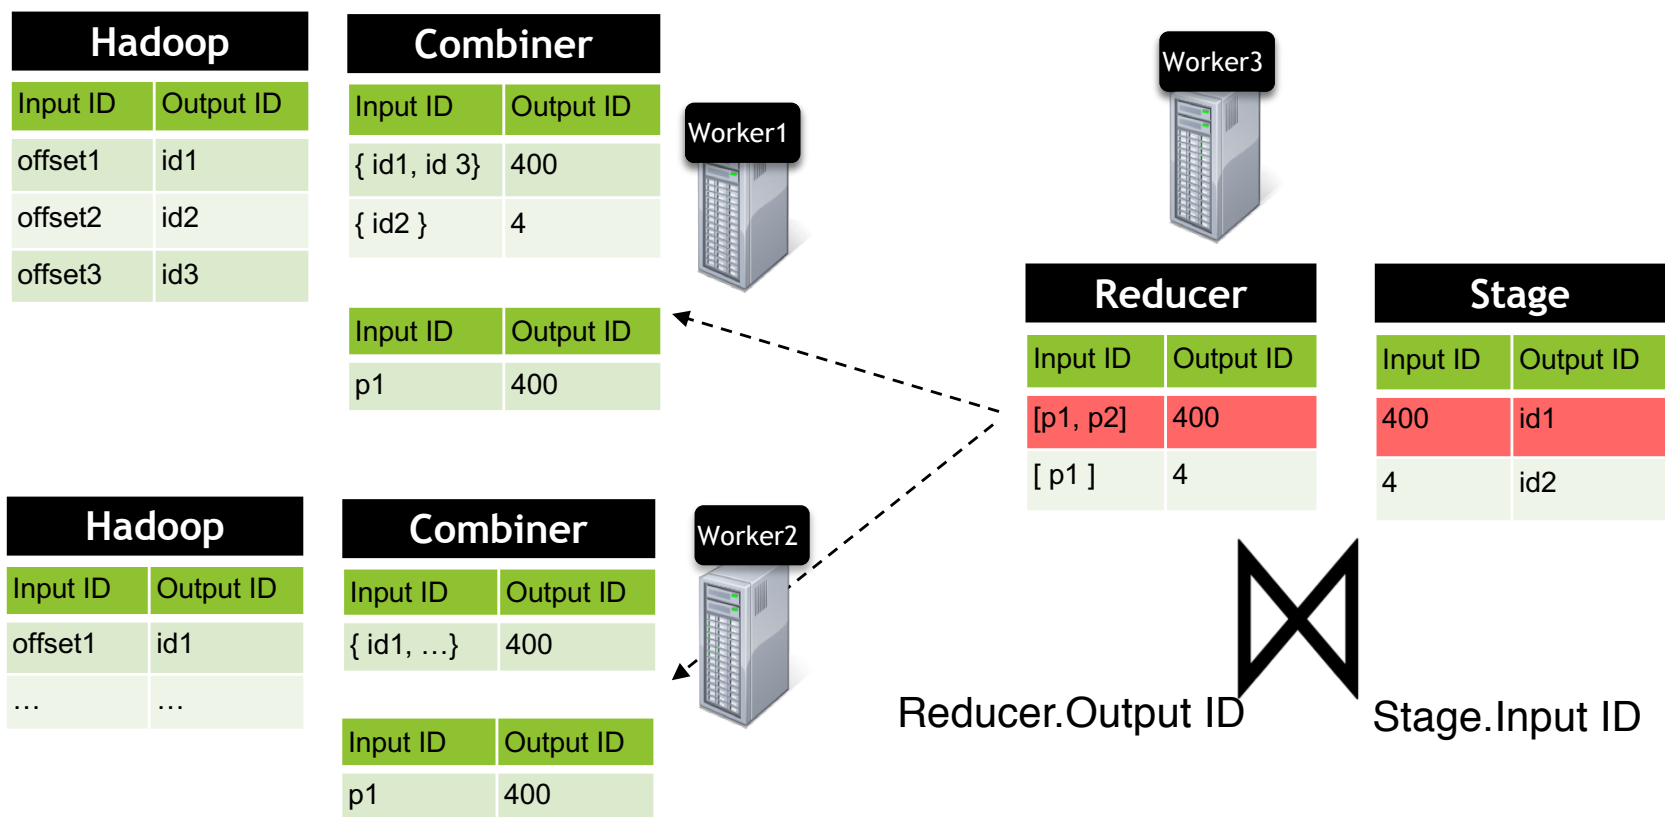
Sensors			
Tuple-ID	Time	Sender-ID	Temperature
T1	11AM	1	34
T2	11AM	2	35
T3	11AM	3	35
T4	12PM	1	35
T5	12PM	2	35
T6	12PM	3	100
T7	1PM	1	35
T8	1PM	2	35
T9	1PM	3	80

Outlier
Outlier

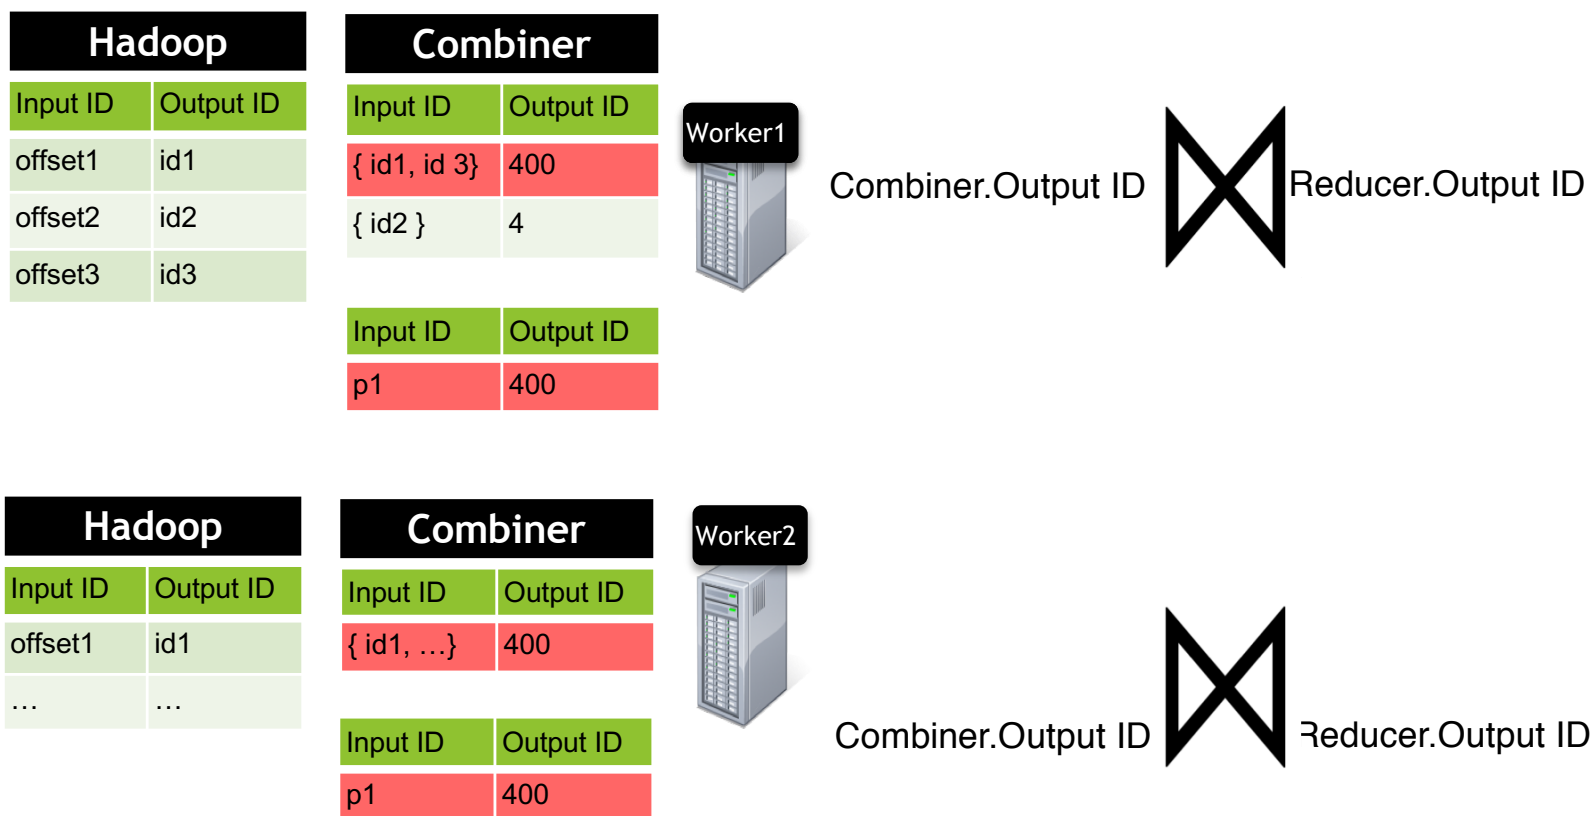
Step 1: Instrument Workflow in Spark



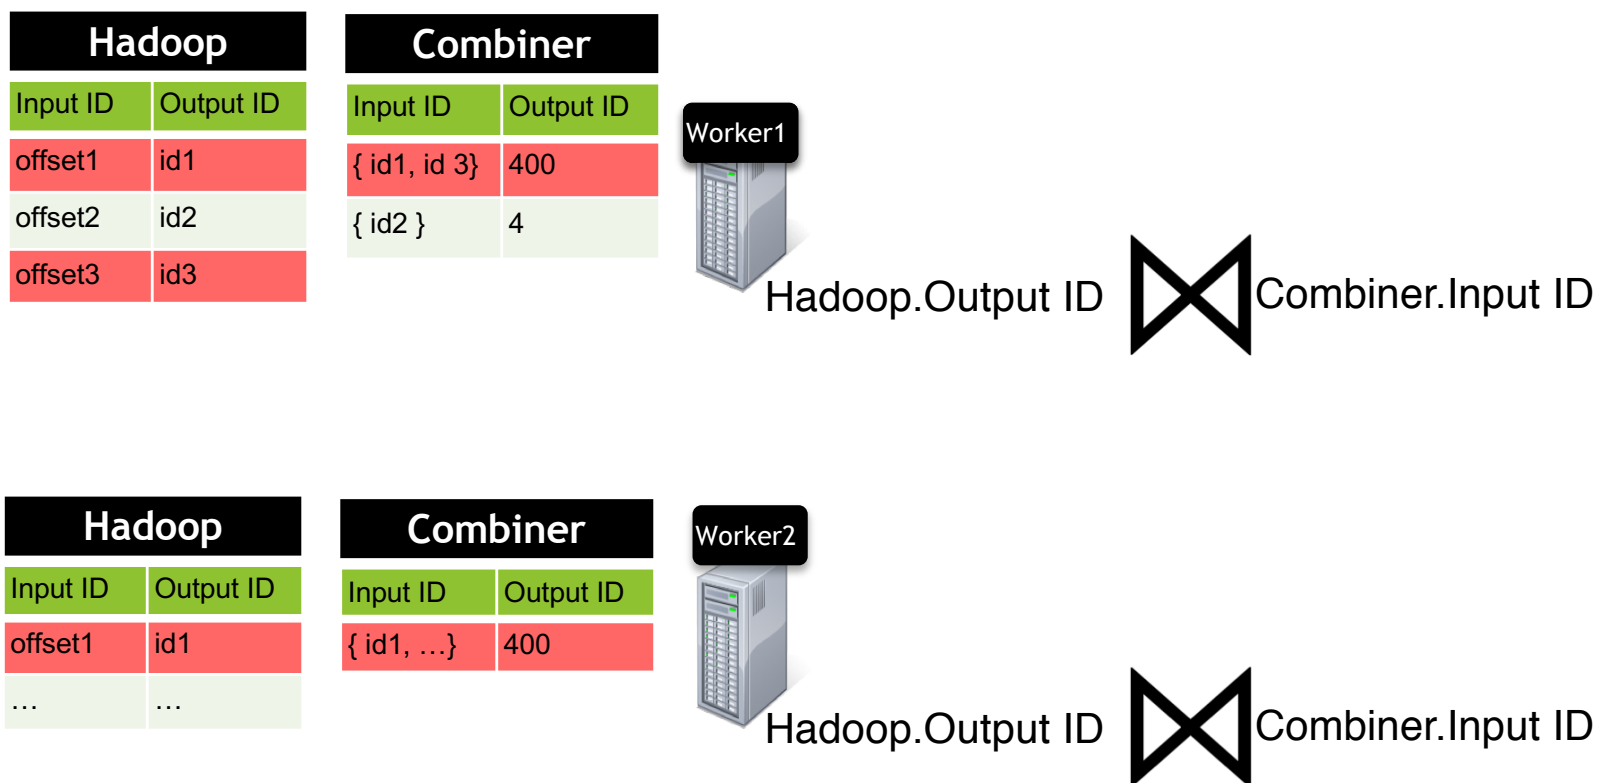
Step 2: Example Backward Tracing



Step 2: Example Backward Tracing



Step 2: Example Backward Tracing



BigSift: Automated Debugging of Big Data Analytics in DISC Applications

SoCC 2017

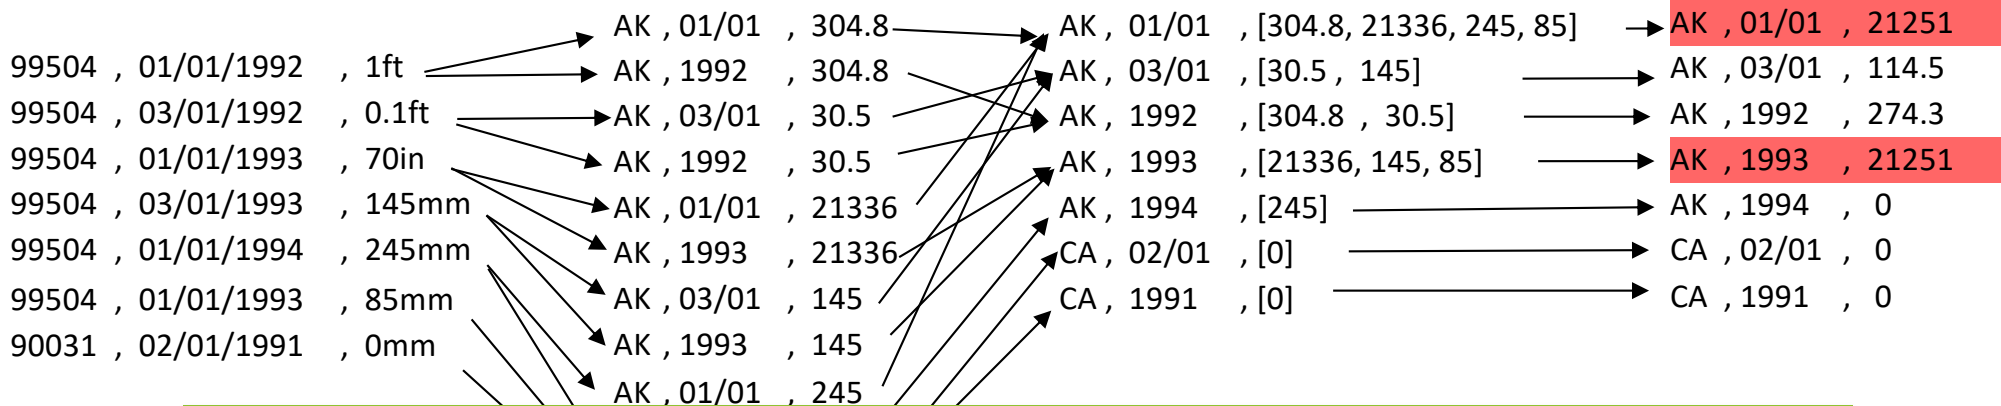
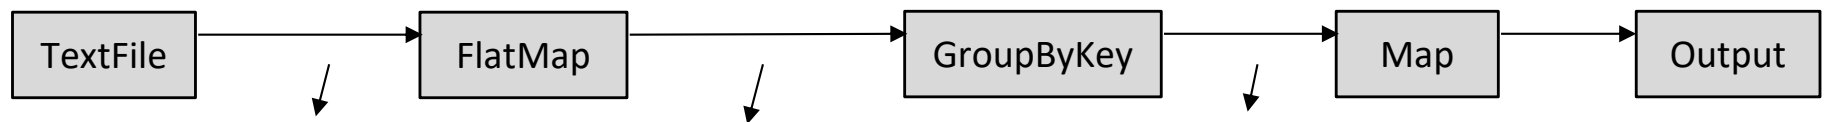
Muhammad Ali Gulzar, Matteo Interlandi, Xueyuan Han, Tyson Condie, Miryung Kim

Motivating Example for Automated Debugging with BigSift

- Alice writes a Spark program that identifies, **for each state** in the US, the **delta between the minimum and the maximum** snowfall reading for **each day of any year** and **for any particular year**.
- An input data record that measures 1 foot of snowfall on January 1st of Year 1992, in the 99504 zip code (Anchorage, AK) area, appears as

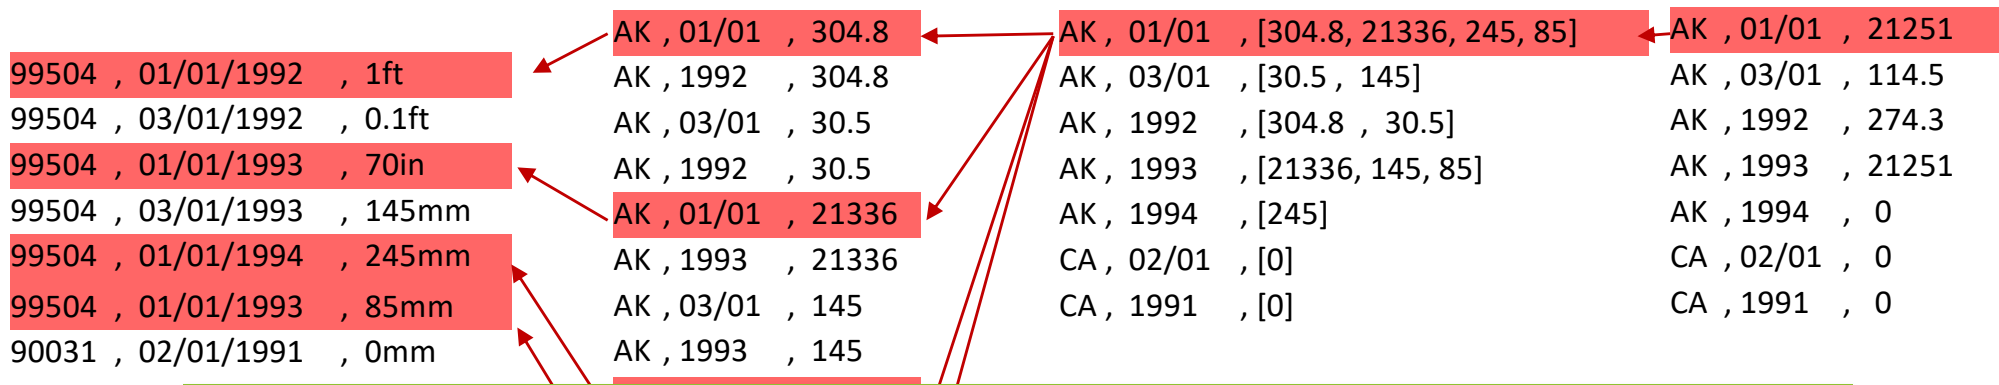
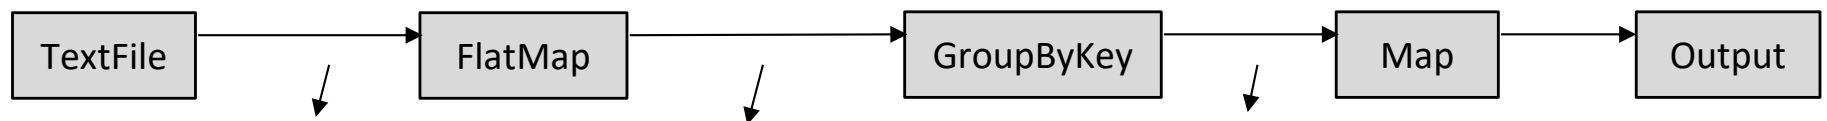
99504 , 01/01/1992 , 1ft

Debugging Incorrect Output



It is challenging because the input records is very large and the program involves computing min and max and a unit conversion

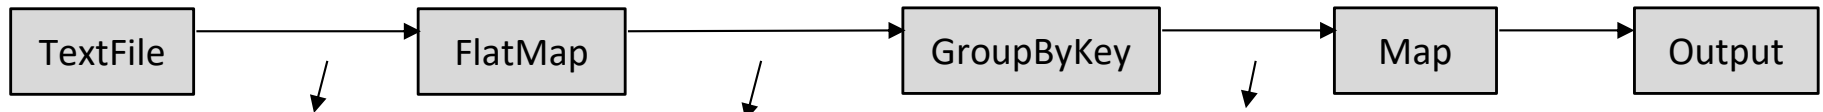
Data Provenance with Titian [VLDB '15]



It over-approximates the scope of failure-inducing inputs *i.e.* records in the faulty key-group are all marked as faulty

Delta Debugging

- Alice tests the output from different input configurations using a test function



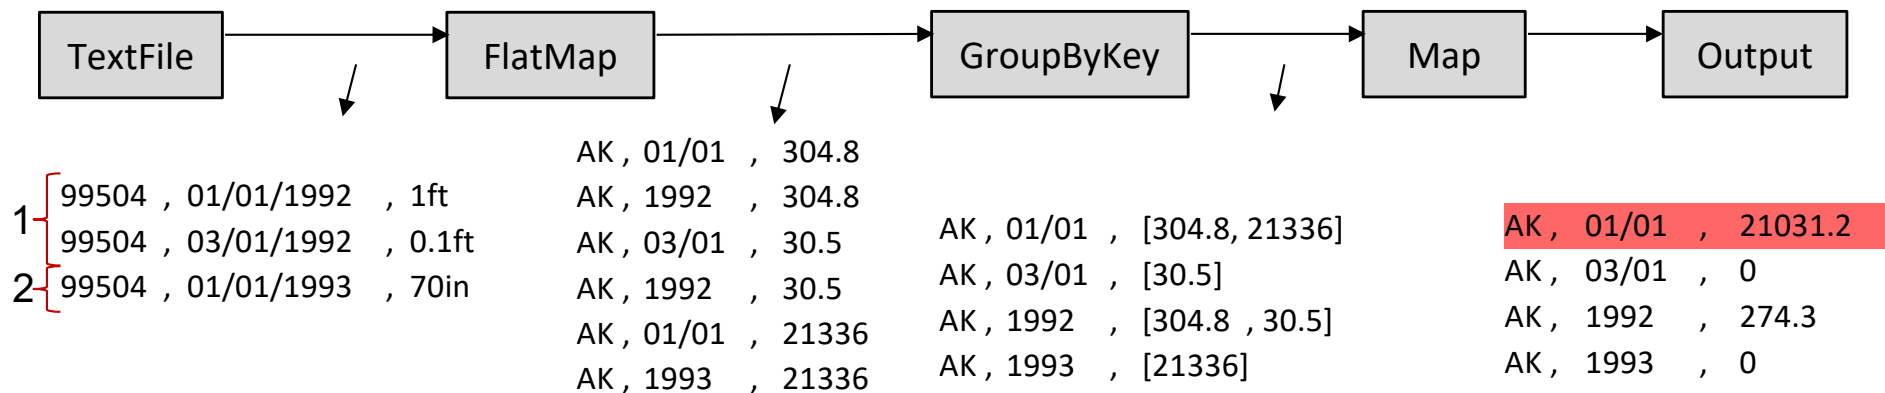
```
def test(key:String, delta: Float) : Boolean = {
  delta < 6000
}
```

1	99504 , 01/01/1992						
	99504 , 03/01/1992						
2	99504 , 01/01/1993 , 70in	AK , 1993 , 21336	AK , 03/01 , [30.5 , 145]				
	99504 , 03/01/1993 , 145mm	AK , 03/01 , 145	AK , 1992 , [304.8 , 30.5]				
	99504 , 01/01/1994 , 245mm	AK , 1993 , 145	AK , 1993 , [21336, 145, 85]				
	99504 , 01/01/1993 , 85mm	AK , 01/01 , 245	AK , 1994 , [245]				
	90031 , 02/01/1991 , 0mm	AK , 1994 , 245	CA , 02/01 , [0]				
			CA , 1991 , [0]				

Run 1 **fails** the test

Delta Debugging

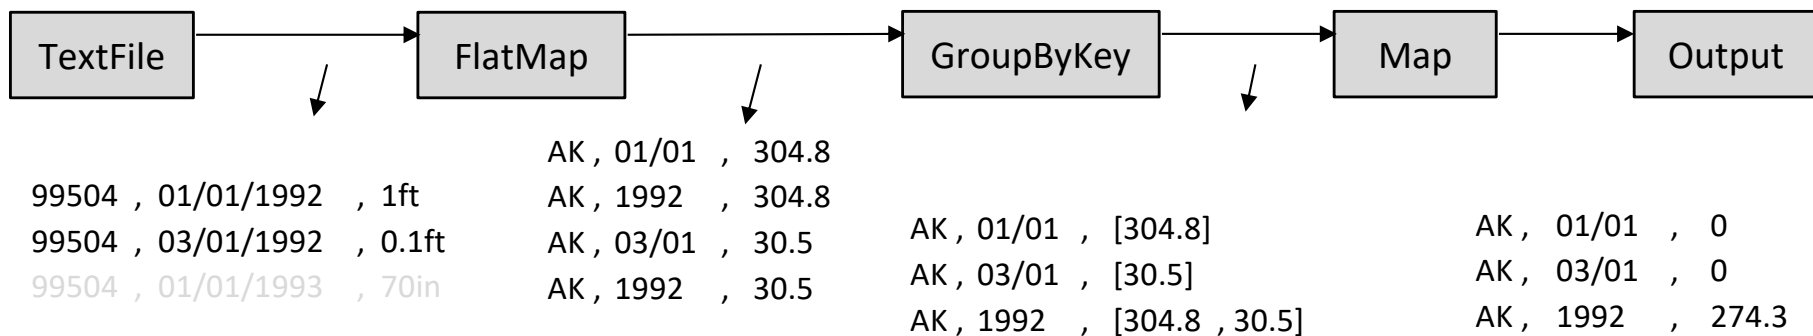
- DD is an systematic debugging procedure guided by a test function.
DD is **inefficient** because it does not eliminate irrelevant input records upfront.



Run 2 **fails** the test

Delta Debugging

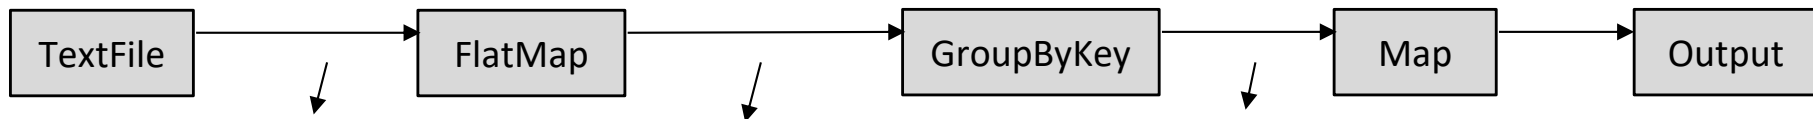
- DD is an systematic debugging procedure guided by a test function.
DD is **inefficient** because it does not eliminate irrelevant input records upfront.



Run 3 **passes** the test

Delta Debugging

- DD is an systematic debugging procedure guided by a test function.
DD is **inefficient** because it does not eliminate irrelevant input records upfront.

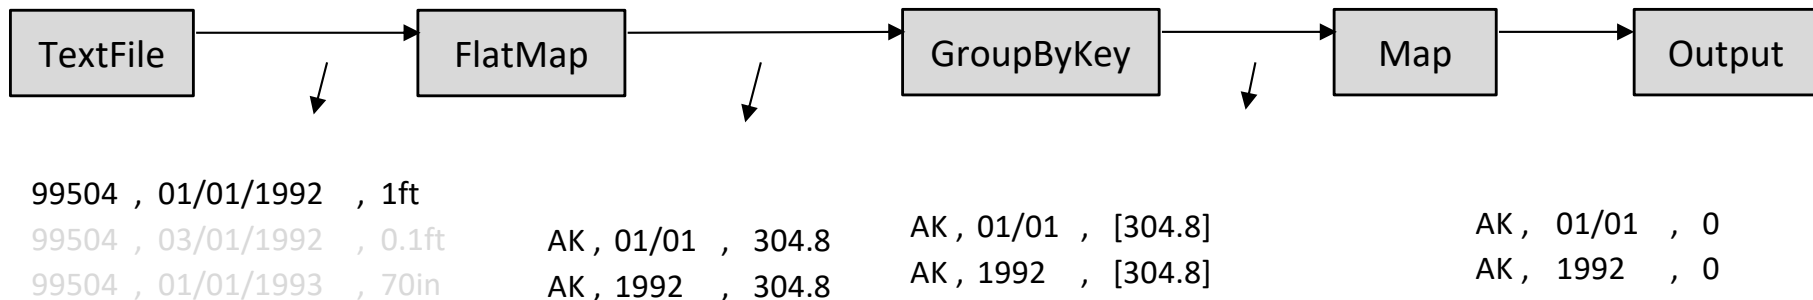


99504 , 01/01/1992 , 1ft			
99504 , 03/01/1992 , 0.1ft	AK , 01/01 , 21336	AK , 01/01 , [21336]	AK , 01/01 , 0
99504 , 01/01/1993 , 70in	AK , 1993 , 21336	AK , 1993 , [21336]	AK , 1993 , 0

Run 4 **passes** the test

Delta Debugging

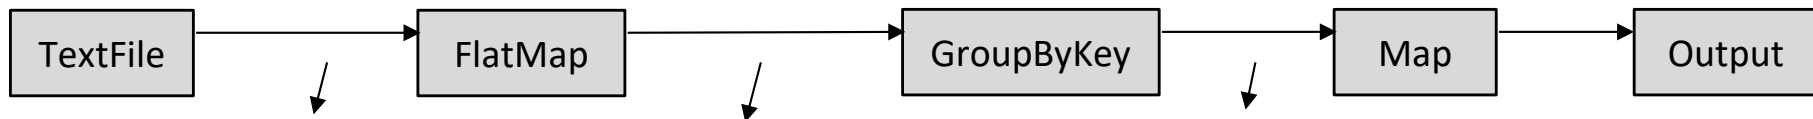
- DD is an systematic debugging procedure guided by a test function.
DD is **inefficient** because it does not eliminate irrelevant input records upfront.



Run 5 **passes** the test

Delta Debugging

- DD is an systematic debugging procedure guided by a test function.
DD is **inefficient** because it does not eliminate irrelevant input records upfront.

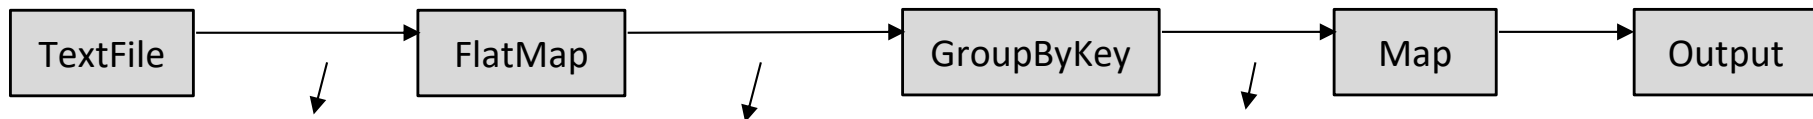


99504 , 01/01/1992 , 1ft			
99504 , 03/01/1992 , 0.1ft	AK , 03/01 , 30.5	AK , 03/01 , [30.5]	AK , 03/01 , 0
99504 , 01/01/1993 , 70in	AK , 1992 , 30.5	AK , 1992 , [30.5]	AK , 1992 , 0

Run 6 **passes** the test

Delta Debugging

- DD is a systematic debugging procedure guided by a test function. DD is **inefficient** because it does not eliminate irrelevant input records upfront.

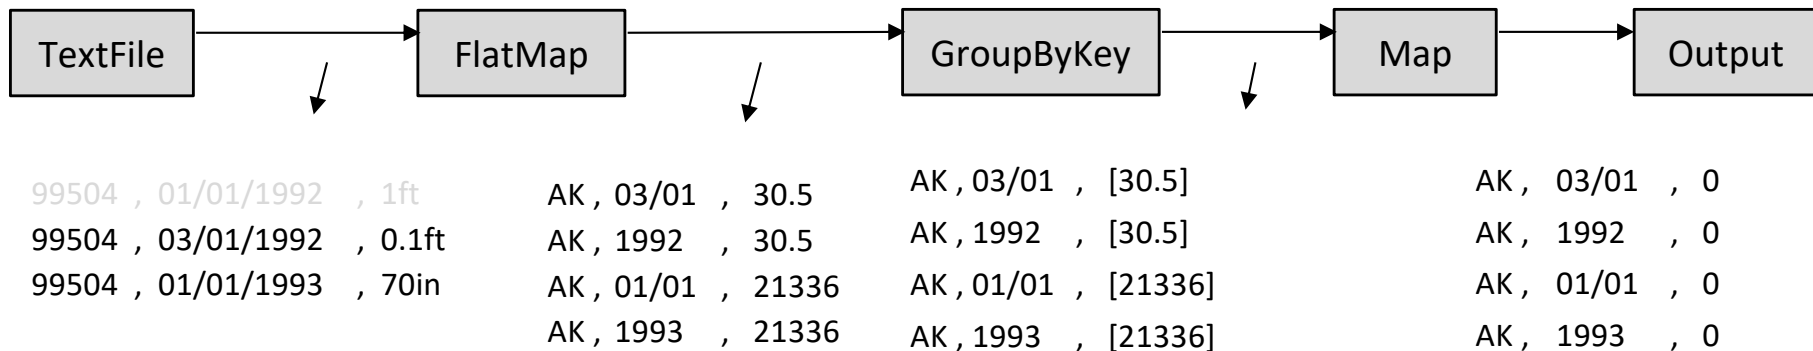


99504 , 01/01/1992 , 1ft			
99504 , 03/01/1992 , 0.1ft	AK , 01/01 , 21336	AK , 01/01 , [21336]	AK , 01/01 , 0
99504 , 01/01/1993 , 70in	AK , 1993 , 21336	AK , 1993 , [21336]	AK , 1993 , 0

Run 7 **passes** the test

Delta Debugging

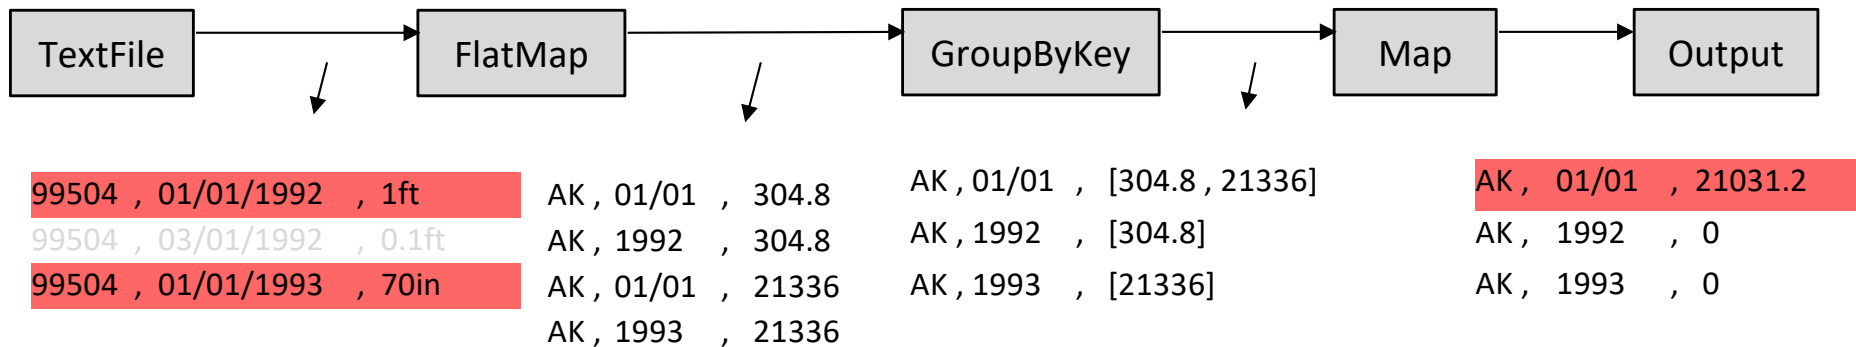
- DD is an systematic debugging procedure guided by a test function.
DD is **inefficient** because it does not eliminate irrelevant input records upfront.



Run 8 **passes** the test

Delta Debugging

- DD is a systematic debugging procedure guided by a test function. DD is **inefficient** because it does not eliminate irrelevant input records upfront.



Run 9 **fails** the test

Automated Debugging in DISC with BigSift

Input: A Spark Program, A Test Function



Output: Minimum Fault-Inducing Input Records



Data Provenance + Delta Debugging

Test Predicate Pushdown

Prioritizing Backward Traces

Bitmap based Test Memoization

Optimization 1: Test Predicate Pushdown

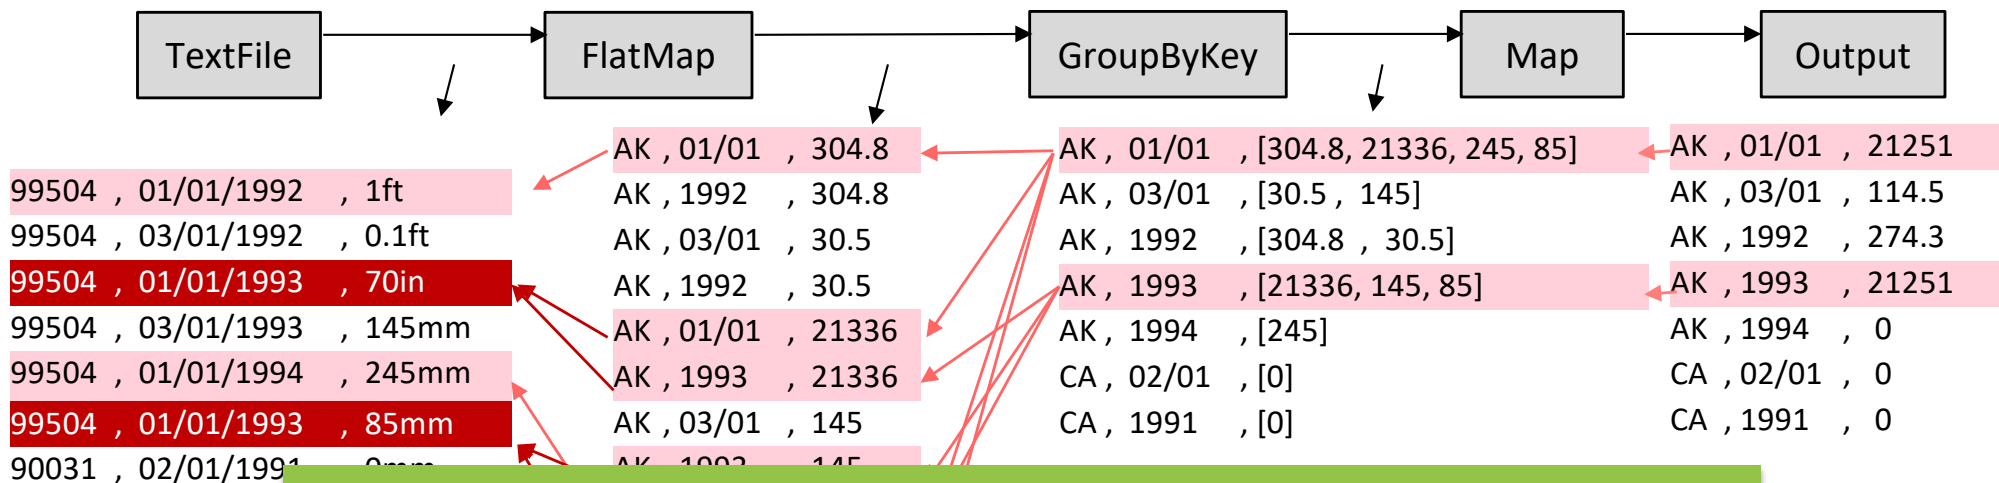
- Observation:** During backward tracing, data provenance traces through all the partitions even though only a few partitions are faulty



If applicable, BigSift pushes down the test function to test the output of combiners in order to isolate the faulty partitions.

Optimization 2 :Prioritizing Backward Traces

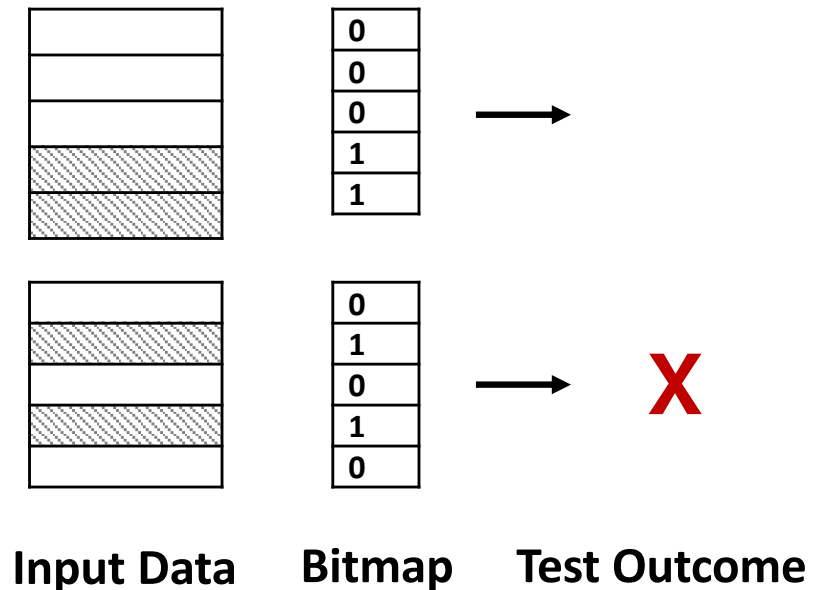
- Observation:** The same faulty input record may contribute to multiple output records failing the test.



In case of multiple faulty outputs, BigSift overlaps two backward traces to minimize the scope of fault-inducing input records

Optimization 3: Bitmap based Test Memoization

- **Observation:** Delta debugging may try running a program on the same subset of input redundantly.
- BigSift leverages bitmap to compactly encode the offsets of original input to refer to an input subset



We use a bitmap based test memoization technique to avoid redundant testing of the same input dataset.



SPARK+AI
SUMMIT 2018

Demo

Debugging Time

Subject Program		Running Time (sec)	Debugging Time (sec)		
Subject Program	Fault	Original Job	DD	BigSift	Improvement
Movie Histogram	Code	56.2	232.8	17.3	13.5X
Inverted Index	Code	107.7	584.2	13.4	43.6X
Rating Histogram	Code	40.3	263.4	16.6	15.9X
Sequence Count	Code	356.0	13772.1	208.8	66.0X
Rating Frequency	Code	77.5	437.9	14.9	29.5X
College Student	Data	53.1	235.3	31.8	7.4X
Weather Analysis	Data	238.5	999.1	89.9	11.1X
Transit Analysis	Code	45.5	375.8	20.2	18.6X

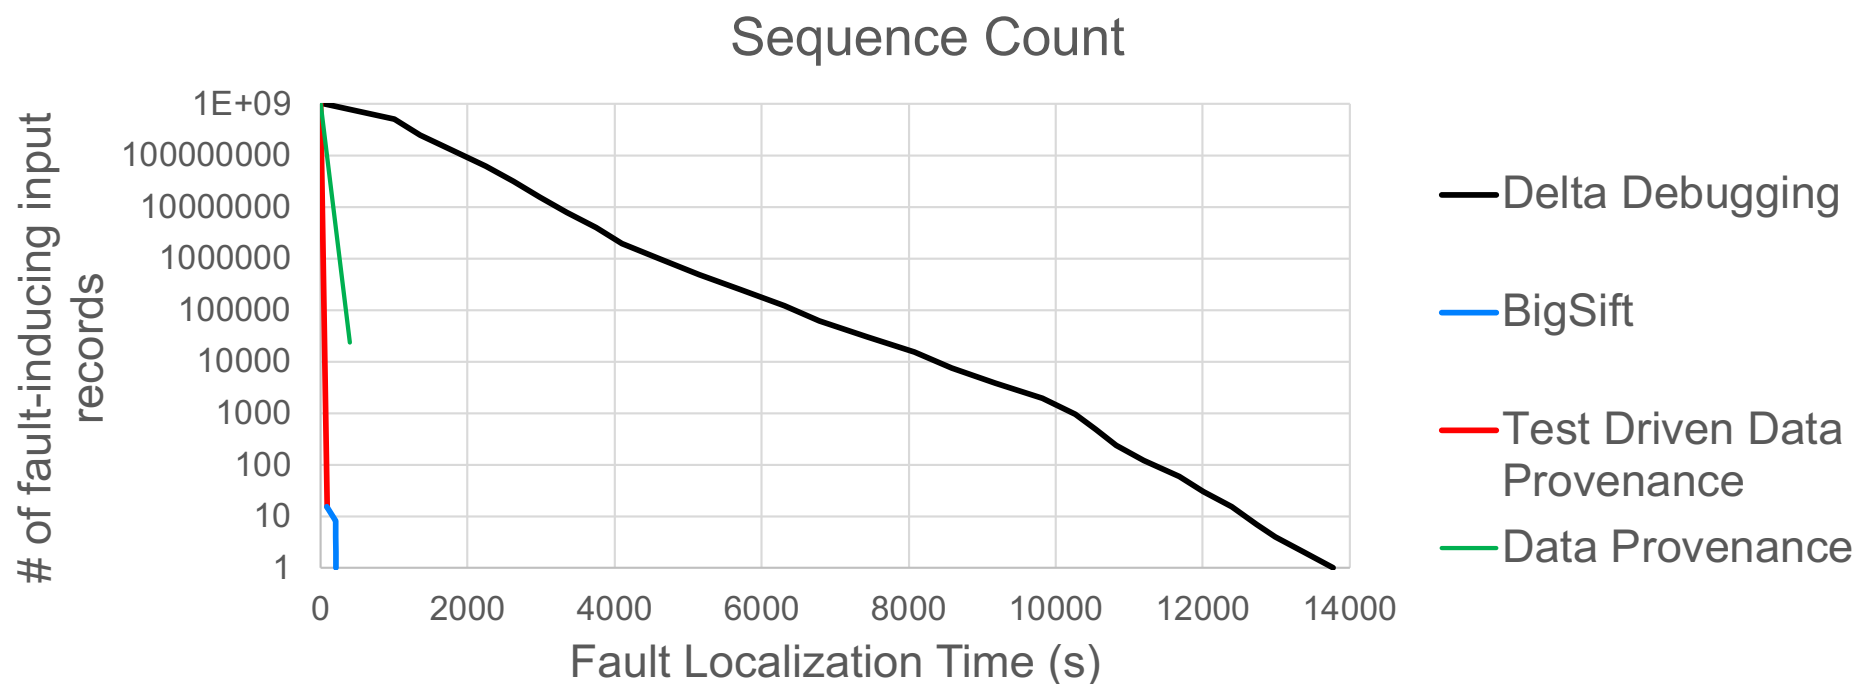
BigSift provides up to a 66X speed up in isolating the precise fault-inducing input records, in comparison to the baseline DD

Debugging Time

Subject Program		Running Time (sec)	Debugging Time (sec)		
Subject Program	Fault	Original Job	DD	BigSift	Improvement
Movie Histogram	Code	56.2	232.8	17.3	13.5X
Inverted Index	Code	107.7	584.2	13.4	43.6X
Rating Histogram	Code	40.3	263.4	16.6	15.9X
Sequence Count	Code	356.0	13772.1	208.8	66.0X
Rating Frequency	Code	77.5	437.9	14.9	29.5X
College Student	Data	53.1	235.3	31.8	7.4X
Weather Analysis	Data	238.5	999.1	89.9	11.1X
Transit Analysis	Code	45.5	375.8	20.2	18.6X

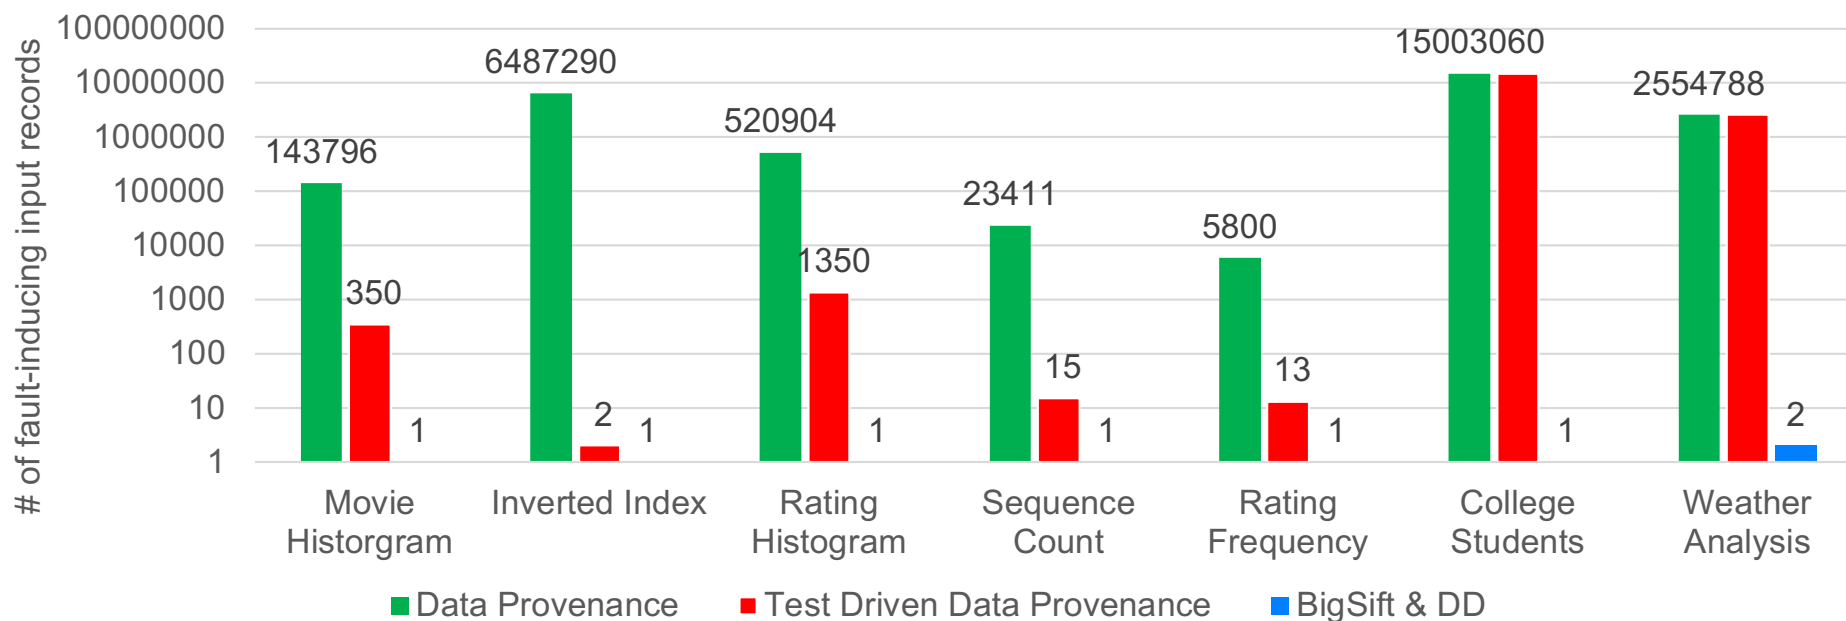
On average, BigSift takes 62% less time to debug a single faulty output than the time taken for a single run on the entire data.

Debugging Time



On average, BigSift takes 62% less time to debug a single faulty output than the time taken for a single run on the entire data.

Fault Localizability over Data Provenance



BigSift leverages DD after DP to continue fault isolation, achieving several orders of magnitude 10^3 to 10^7 better precision.

Summary

- BigSift is the first piece of work in automated debugging of big data analytics in DISC.
- BigSift provides **$10^3X - 10^7X$ more precision** than data provenance in terms of fault localizability.
- It provides up to **66X speed up** in debugging time over baseline Delta Debugging.
- In our evaluation we have observed that, on average, BigSift finds the faulty input in **62% less** than the original job execution time.