

# Cloud Cost Management and Apache Spark

Xuan Wang, Databricks

**#DSSAIS13**

# Introduction

- Goal of this talk
  - share our experience in managing cloud costs
  - tools and technologies
  - lessons learnt and good practices
  - go wide rather than go deep

# Introduction

- Goal of this talk
  - share our experience in managing cloud costs
  - tools and technologies
  - lessons learnt and good practices
  - go wide rather than go deep
- Why do we care about cloud cost?
  - growth in cloud revenue in Q1 2018: Amazon: 49%, Microsoft: 58%

# Databricks' Unified Analytics Platform

Unifies Data  
Engineers and Data  
Scientists

Unifies Data and  
AI Technologies

Eliminates  
infrastructure  
complexity

## COLLABORATIVE NOTEBOOKS



Data Engineers



Data Scientists

## DATABRICKS RUNTIME

Powered by 

Delta

SQL

Streaming

XGBoost



CLOUD NATIVE SERVICE



# Three paths toward cost control

- Native reporting from cloud providers
  - Good general information and supports
  - Limited options, not scalable as environment grows
- Commercial tools
  - More details and flexibilities, connectors to raw data
  - Not enough customization, additional charges

# Three paths toward cost control

- Native reporting from cloud providers
  - Good general information and supports
  - Limited options, not scalable as environment grows
- Commercial tools
  - More details and flexibilities, connectors to raw data
  - Not enough customization, additional charges
- In-house solutions
  - **Most flexible, deeper understanding of the costs**
  - Opportunity costs


# Challenges in cloud cost control

- overwhelming and complex usage details
  - need to convert data into insights/actions
- gaps between “hands” and “wallets”
  - developers consume resources without realizing the charges
- evolving cloud landscape
  - external: new services, new discounts, ...
  - internal: new use cases, new architecture, ...

# Our solutions

## Raw Data

cost and usage  
s3 access logs  
s3 inventory  
ec2/rds snapshot  
reserved instances  
...



## Analytics

Databricks  
Notebooks  
  
BI tools:  
Superset, Tableau,  
...  
  
Monitors and  
alerts



# Our solutions

## Raw Data

cost and usage  
s3 access logs  
s3 inventory  
ec2/rds snapshot  
reserved instances  
...

The **data** problem:  
ETL and attribute costs

**DATABRICKS  
DELTA**

**DATA LAKE**

## Analytics

Databricks  
Notebooks

BI tools:  
Superset, Tableau,  
...

Monitors and  
alerts

The **process** problem:  
prioritize, optimize, monitor,  
automate

# The data problem

- cost and usage report (detailed billing)
  - CSV, grouped by month, updated daily

# The data problem

- cost and usage report (detailed billing)
  - CSV, grouped by month, updated daily
- EC2/RDS snapshots and reserved instances
  - JSON, from REST API


# The data problem

- cost and usage report (detailed billing)
  - CSV, grouped by month, updated daily
- EC2/RDS snapshots and reserved instances
  - JSON, from REST API
- S3 inventory
  - CSV/ORC, snapshot, updated daily/weekly
- S3 access logs
  - raw logs in text, updated multiple times a day

# The data problem

- cost and usage report (detailed billing)
  - **CSV**, grouped by month, updated **daily**
- EC2/RDS snapshots and reserved instances
  - **JSON**, from **REST API**
- S3 inventory
  - **CSV/ORC**, snapshot, updated **daily/weekly**
- S3 access logs
  - **raw logs in text**, updated **multiple times a day**

# Data pipelines with Spark



## Challenges

- Data corruptions
- Multiple jobs/staging tables
- Reliability and consistency

# Databricks Delta: Analytics Ready Data

## 1. Data Reliability

ACID Compliant Transactions  
Schema Enforcement & Evolution

## 2. Query Performance

Very Fast at Scale  
Indexing & Caching (10-100x Faster)

### LOTS OF NEW DATA

Customer Data  
Click Streams  
Sensor data (IoT)  
Video/Speech  
...



## 3. Simplified Architecture

Unify batch & streaming  
Early data availability for analytics

# ETL: AWS cost and usage

```
spark.read
  .option("header", "true")
  .csv(inputFiles: _*)
  .write
  .format("delta")
  .option("replaceWhere", s"month = '$month'")
  .save(outputDir)
```



# ETL: AWS cost and usage

```
spark.read
  .option("header", "true")
  .csv(inputFiles: _*)
  .write
    .format("delta")
    .option("replaceWhere", s"month = '$month'")
  .save(outputDir)
```

# ETL: AWS s3 access logs

```
// Fake s3 access logs
""""
3E57427F33A59F07 [06/Feb/2014:00:00:38 +0000] REST.PUT.OBJECT "GET /myobj HTTP/1.1" 200 NoSuchBucket"
""""
```

```
def parseS3AccessLogRecord(s: String): S3AccessLogRecord = ???
```

```
spark.read
```

```
  .textFile(inputFiles:_) // Dataset[String]
```

```
  .map(parseS3AccessLogRecord) // Dataset[S3AccessLogRecord]
```

```
  .write
```

```
  .format("delta")
```

```
  .option("replaceWhere", s"date >= '$startDate' AND date < '$endDate'")
```

```
  .save(outputDir)
```

# Manage Databricks Delta tables

- Create table

```
CREATE TABLE s3_access_logs USING delta LOCATION '$path'
```

- Optimize table

```
OPTIMIZE s3_access_logs ZORDER BY bucket
```

# Manage Databricks Delta tables

- Create table

```
CREATE TABLE s3_access_logs USING delta LOCATION '$path'
```


- Optimize table

```
OPTIMIZE s3_access_logs ZORDER BY bucket
```

- Query table

```
SELECT * FROM s3_access_logs WHERE bucket = 'my-bucket'
```

**Files layout &  
statistics:**



Delta Logs:  
File1: min='a', max='g'  
File2: min='g', max='n'  
File3: min='o', max='z'

# Attributions


- Rule based attributions
  - accounts
    - dedicated accounts for different teams / use cases
  - tagging
    - tag resources with budget groups
  - manual rules
    - should avoid this as much as possible

# The process problem

- **Prioritize**
  - high data transfer cost
- **Optimize**
  - reserved instance purchases
- **Monitor**
  - predictions and alerts
- **Automate**
  - auto-shutdown unused resources


# Story: high data transfer cost

- Observation
  - Cross region data transfers are expensive
  - Two buckets cost about \$1k/day



# Story: high data transfer cost


- Observation
  - Cross region data transfers are expensive
  - Two buckets cost about \$1k/day
- Root cause
  - downloading spark images






# Story: high data transfer cost

- Actions
  - Distribute images to multiple regions.
  - Monitor on cross region cost



# Story: high data transfer cost

- Actions
  - Distribute images to multiple regions.
  - Monitor on cross region cost
- Results
  - Significantly reduced cost
  - Faster cluster creation



# Optimization: reserved instances

- Reserved instances (RI)
  - 1-yr/3-yr commitment in exchange for discounts
  - underutilized instances, upfront cost
  - significant discounts, availability

# Optimization: reserved instances

- Reserved instances (RI)
  - 1-yr/3-yr commitment in exchange for discounts
  - underutilized instances, upfront cost
  - significant discounts, availability
- Challenges
  - non-trivial to decide how much RI to purchase
  - need to predict the future

# Optimization: reserved instances

- Assign budgets to teams
- Provide tool to compute the optimal RI to buy
- Define process for RI purchase requests and approvals

0. AWS account ID :

1. AWS product :

2. start date (yyyy-mm-dd) :

3. end date (exclusive) :

4. min RI utilization :

5. custom SQL filter :

### Plan EC2/RDS RI purchase (go/ri/planner)

This notebook helps you estimate how many EC2/RDS RIs you should purchase, in which AWS account and region. The estimation is based on past on-demand usage during a specific period of time. And we assume the usage will be the same for at least one year. The RI quota is shared within each instance family (e.g., m4) and region, across all accounts under AWS organization. For example, buying 40 m4.xlarge RIs in us-west-2 is equivalent to buying 10 m4.4xlarge RIs in the same region.


Parameters:

- **accountId:** AWS account ID. You can look up account ID from `aws.accounts` table. If not specified, the notebook will check all AWS accounts.
- **product:** AWS product, either EC2 or RDS
- **startDate:** start date of the duration
- **endDate:** end date (exclusive) of the duration. It must be at least three days ago to handle late record.
- **minRiUtilization:** minimal RI utilization, default 0.9

# Monitor and alerts

- Why
  - prevent degenerations
  - proactive to “bill shock”
- Challenges
  - different patterns for different use cases
  - changing baselines


# Monitor and alerts



Picture from Sharma, S., Swayne, D.A. & Obimbo, C. *Energy, Ecology and Environment*. (2016) 1: 123.  
<https://doi.org/10.1007/s40974-016-0011-1>

# Monitor and alerts

- adaptive prediction with change point detection
- alerts for each budget group
- scheduled jobs and dashboards





# Auto-shutdown with Custodian

- <https://github.com/capitalone/cloud-custodian>
  - Rule based cloud infrastructure management tool



# Auto-shutdown with Custodian

```
1  policies:
2  - actions:
3    - terminate
4    comment: Terminate dev EC2 instances that are more than 10 hours old
5    filters:
6      - tag:Owner: absent
7      - tag:Project: absent
8      - State.Name: running
9      - key: tag:Name
10      op: regex
11      type: value
12      value: ^development.*$
13      - hours: 10
14      type: instance-age
15      name: terminated-leaked-dev-ec2
16      resource: ec2
```

# Auto-shutdown with Custodian

```
1  policies:
2  - actions:
3    - terminate
4    comment: Terminate dev EC2 instances that are more than 10 hours old
5    filters:
6      - tag:Owner: absent
7      - tag:Project: absent
8      - State.Name: running
9      - key: tag:Name
10        op: regex
11        type: value
12        value: ^development.*$
13      - hours: 10
14        type: instance-age
15    name: terminated-leaked-dev-ec2
16    resource: ec2
```

# Auto-shutdown with Custodian

```
1  policies:
2  - actions:
3    - terminate
4    comment: Terminate dev EC2 instances that are more than 10 hours old
5    filters:
6    - tag:Owner: absent
7    - tag:Project: absent
8    - State.Name: running
9    - key: tag:Name
10      op: regex
11      type: value
12      value: ^development.*$
13    - hours: 10
14      type: instance-age
15    name: terminated-leaked-dev-ec2
16    resource: ec2
```

# Summary

- in-house solutions because
  - flexibility and deeper understanding of cost and usage
- cost attribution - a **data** problem
  - ETL: Databricks Delta for analytics ready data
  - explore: Databricks Notebooks and BI tools via JDBC
  - attribute: rule-based, tagging is important
- cost control - a **process** problem
  - prioritize: get the work done!
  - optimize: distributed ownerships, and centralized tools
  - monitor: change points + basic linear model
  - automate: Custodian for managing cloud infrastructure

# Thank you

[xwang@databricks.com](mailto:xwang@databricks.com)

<https://www.linkedin.com/in/xuanwang2>

<https://databricks.com/product/unified-analytics-platform>

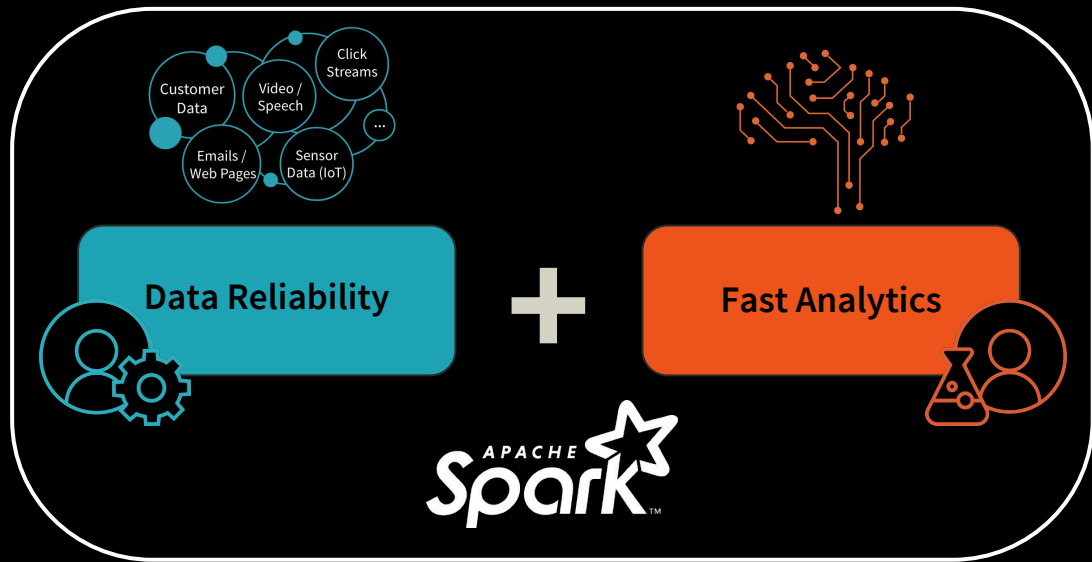
# BACKUP SLIDES

Xuan Wang, Databricks

**#AssignedHashtagGoesHere**

# New: Databricks Delta

*Extends Apache Spark to simplify data **reliability** and **performance***





# The data problem

- cost and usage report
  - detailed usage and billing information by hours
  - CSV, delivered at least once a day
- s3 inventory
  - list of all objects in s3 and associated metadata
  - CSV/ORC, delivered daily/weekly

Bucket	Key	VersionId	IsLatest	IsDeleteMarker	Size	LastModifiedDate	Etag	StorageClass	MultipartUploaded	ReplicationStatus
example-bucket	object1			FALSE	2.4E+08	2016-08-11T01:19	e80d8eda4	STANDARD	TRUE	
example-bucket	object2			FALSE	0	2016-08-10T22:23	d41d8cd98	STANDARD	FALSE	
example-bucket	object3			FALSE	9	2016-08-10T20:18	9090441e4	STANDARD_IA	FALSE	
example-bucket	object4			FALSE	9	2016-08-10T20:36	9090441e4	STANDARD_IA	FALSE	
example-bucket	object1			FALSE	22	2016-08-10T20:35	9090441e4	STANDARD	FALSE	
example-bucket	object1			FALSE		2016-08-10T20:34	9090441e4	REDUCED_RED	FALSE	
example-bucket	object1			FALSE		2016-08-10T21:13	9090441e4	GLACIER	FALSE	

# The data problem

- s3 access logs
  - requests/API calls to access s3 objects
  - raw logs, delivered frequently

```
// Fake s3 access logs
*****
3E57427F33A59F07 [06/Feb/2014:00:00:38 +0000] REST.PUT.OBJECT "GET /myobj HTTP/1.1" 200 NoSuchBucket"
*****
```

- EC2/RDS snapshots and reserved instances
  - json from REST API