

Extending Spark SQL API with Easier to Use Array Types Operations

Marek Novotny, ABSA Capital

Jan Scherbaum, ABSA Capital

#Dev3SAIS

Intro



- ABSA (Barclays Africa) is a Pan-African financial services provider
 - With Apache Spark at the core of its data engineering
- We try to fill gaps in the Hadoop eco-system
 - Especially from the perspective of highly regulated industries
 - ABRiS – Avro SerDe for Apache Spark structured APIs
 - <https://github.com/AbsaOSS/ABRiS>
 - Spline – a lineage tracking for Apache Spark
 - <https://absaoss.github.io/spline/>

Problem Statement

- Lots of structured data (XML, JSON) from a data lake has to be cleansed and transformed
 - With lineage tracking via Spline
- Limited support for nested structures in Spark structured APIs
- Alternatives
 - RDD/Dataset higher-order functions & UDFs \Rightarrow No lineage tracking
 - Flatten structure with explode \Rightarrow Very inefficient
 - 1000 records, 2 levels, 150 elements per level \Rightarrow 22,500,000 records
 - Convert array elements to columns \Rightarrow too big logical plan
 - Size of the logical plan is directly proportional to max size of arrays
 - Variable array size

Solution

- Create a custom build of Spark with an extension
 - **concat, flatten, reverse**, zip_with_index, transform
 - github.com/AbsaOSS/spark/tree/branch-2.2.1-ABSA
 - github.com/AbsaOSS/spark/tree/branch-2.3.0-ABSA
- Contribute some array functions to Spark 😊
- Apache Spark has gotten inspired by Presto
 - 33 array/map functions are planned for the version 2.4.0
 - [SPARK-23899](https://issues.apache.org/jira/browse/SPARK-23899)

concat (col1, col2, ..., colN)

- [SPARK-23736](#) resolved by Marek Novotny
- Adds array types support to the existing “concat” function
 - `concat([1, 2, 3], [4, 5])` → `[1, 2, 3, 4, 5]` `[1, 2, 3] ≈ array(1, 2, 3)`
 - `concat([1, 2], [2, 3], [3, 4])` → `[1, 2, 2, 3, 3, 4]`
- Type coercion
 - `concat([1, 2, 3], ['a', 'b', 'c'])` → `['1', '2', '3', 'a', 'b', 'c']`
- Null handling
 - `concat(..., null, ...)` → `null`
 - `concat([1, null], ['a', null])` → `['1', null, 'a', null]`

flatten (col)

- [SPARK-23821](#) resolved by Marek Novotny

- Performs shallow flattening

<code>flatten([[1, 2, 3], [4, 5]])</code>	→	<code>[1, 2, 3, 4, 5]</code>
<code>flatten([['a', 'b'], [], ['b', 'a']])</code>	→	<code>['a', 'b', 'b', 'a']</code>
<code>flatten([[[1, 2], [3]], [[4]]])</code>	→	<code>[[1, 2], [3], [4]]</code>

- Null handling

<code>flatten(<i>null</i>)</code>	→	<code><i>null</i></code>
<code>flatten([[1, 2, 3], <i>null</i>, [4, 5]])</code>	→	<code><i>null</i></code>
<code>flatten([['a', <i>null</i>], [<i>null</i>, 'b']])</code>	→	<code>['a', <i>null</i>, <i>null</i>, 'b']</code>

reverse (col)

- [SPARK-23926](#) resolved by Marek Novotny
- Adds array types support to the existing “reverse” function
 - `reverse([2, 1, 4, 3, 5])` → `[5, 3, 4, 1, 2]`
 - `reverse(['b', 'a', 'd', 'c'])` → `['c', 'd', 'a', 'b']`
 - `reverse([['a'], [], ['c', 'b'], []])` → `[[], ['c', 'b'], [], ['a']]`
- Null handling
 - `reverse(null)` → `null`
 - `reverse(['a', null, 'b', null])` → `[null, 'b', null, 'a']`

map_entries (col)

- [SPARK-23935](#) resolved by Marek Novotny

- Transforms maps into arrays of key-value pairs

`map_entries(map(2 → 'b', 1 → 'a', 3 → 'c'))` → `[(2, 'b'), (1, 'a'), (3, 'c')]`

`map_entries(map('a' → 1, 'c' → 2, 'b' → 3))` → `[('a', 1), ('c', 2), ('b', 3)]`

- Null handling

`map_entries(null)` → `null`

`map_entries(map(1 → null, 3 → 'c', 2 → null))` → `[(1, null), (3, 'c'), (2, null)]`

map_from_entries (col)

- [SPARK-23934](#) – in progress – github.com/apache/spark/pull/21282
- Transforms arrays of key-value pairs into maps
`map_from_entries([(2, 'a'), (1, 'b'), (3, 'c')])` → `map(2 → 'a', 1 → 'b', 3 → 'c')`
- Null handling
`map_from_entries([(2, 'a'), (1, null)])` → `map(2 → 'a', 1 → null)`
`map_from_entries([(2, 'a'), (null, 'b')])` → **RuntimeException**
`map_from_entries(null)` → `null`

sequence (start, stop[, step])

- [SPARK-23927](#) – in progress – github.com/apache/spark/pull/21155/
- **Integral sequence** (byte, short, int, long)
 - sequence(1, 5) → [1, 2, 3, 4, 5]
 - sequence(1, 5, 2) → [1, 3, 5]
- **Temporal sequence** (date, timestamp)
 - sequence(
 cast('2018-01-01' as date),
 cast('2018-03-01' as date),
 interval 1 month) → [2018-01-01, 2018-02-01, 2018-03-01]
- **Reverse sequence**
 - sequence(5, 1) → [5, 4, 3, 2, 1]
 - sequence(5, 1, -2) → [5, 3, 1]
- Supports time zones and DST

array_max (col)

array_min (col)

- [SPARK-23917](#), [SPARK-23918](#) resolved by Marco Gaido

- Supports any orderable type
(atomic types, structs, arrays, nulls and UDTs)

array_max([2, 5, 3]) → 5

array_min(['b', 'a', 'c']) → 'a'

array_max([[5, 2], [6, 1]]) → [6, 1]

array_min([(5, 2), (6, 1)]) → [5, 2]

array_min([], [0, 2]) → []

- Returns *null* for empty array

array_max([]) → null

array_sort (col)

- [SPARK-23921](#) resolved by Kazuaki Ishizaki
- Sorts arrays in ascending order
 - array_sort([3, 5, 1, 4, 2]) → [1, 2, 3, 4, 5]
 - array_sort(['c', 'a', 'b', 'e', 'd']) → ['a', 'b', 'c', 'd', 'e']
- Nulls handling - *array_sort* vs. *sort_array*
 - array_sort([3, 1, null, 2]) → [1, 2, 3, null]
 - sort_array([3, 1, null, 2], true) → [null, 1, 2, 3]
 - sort_array([3, 1, null, 2], false) → [3, 2, 1, null]

array_join (col, delimiter[, nullRepl])

- [SPARK-23916](#) resolved by Marco Gaido
- Creates a string from an **array of strings** using a delimiter
`array_join(['fish', 'chips'], '&') → 'fish&chips'`
- Can optionally replace nulls
`array_join(['a', null, 'b'], ', ') → 'a, b'`
`array_join(['a', null, 'b'], ', ', '_') → 'a, _, b'`

arrays_overlap (col1, col2)

- [SPARK-23922](#) resolved by Marco Gaido
- Checks for common elements
 - `arrays_overlap([1, 2, 3], [4, 2])` → `true`
 - `arrays_overlap([1, 2, 3], [4, 5])` → `false`
- Null elements are ignored
 - `arrays_overlap([1, null], [1, null])` → **`true`**
- ...but not always
 - `arrays_overlap([1, null], [2, null])` → **`null`**

array_position (col, elem)

element_at (col, index)

- [SPARK-23919](#), [SPARK-23924](#) resolved by Kazuaki Ishizaki
- 1-based indexing like in SQL
- Find a position of the 1st occurrence
 - array_position(['a', 'b'], 'a') → 1
 - array_position(['a', 'b'], 'x') → 0
- Get an element by an index
 - element_at(['a', 'b', 'c'], 2) → 'b'
 - element_at(['a', 'b', 'c'], 42) → *null*
- Get a value by a key
 - element_at(map('name' → 'Bob', 'age' → 42), 'name') → 'Bob'

array_repeat (element, count)

- [SPARK-23925](#) resolved by Florent Pepin

- Repeats anything N times

<code>array_repeat('meh', 3)</code>	→	<code>['meh', 'meh', 'meh']</code>
<code>array_repeat(['doh'], 3)</code>	→	<code>[['doh'], ['doh'], ['doh']]</code>
<code>array_repeat(null, 3)</code>	→	<code>[null, null, null]</code>
<code>array_repeat('foo', 0)</code>	→	<code>[]</code>
<code>array_repeat('bar', -1)</code>	→	<code>[]</code>

slice (col, start, length)

- [SPARK-23930](#) resolved by Marco Gaido
- 1-based indexing
- Returns a sub-array

slice(['a', 'b', 'c', 'd'], 2, 2) → ['b', 'c']

slice(['a', 'b', 'c', 'd'], -2, 2) → ['c', 'd']

In Progress

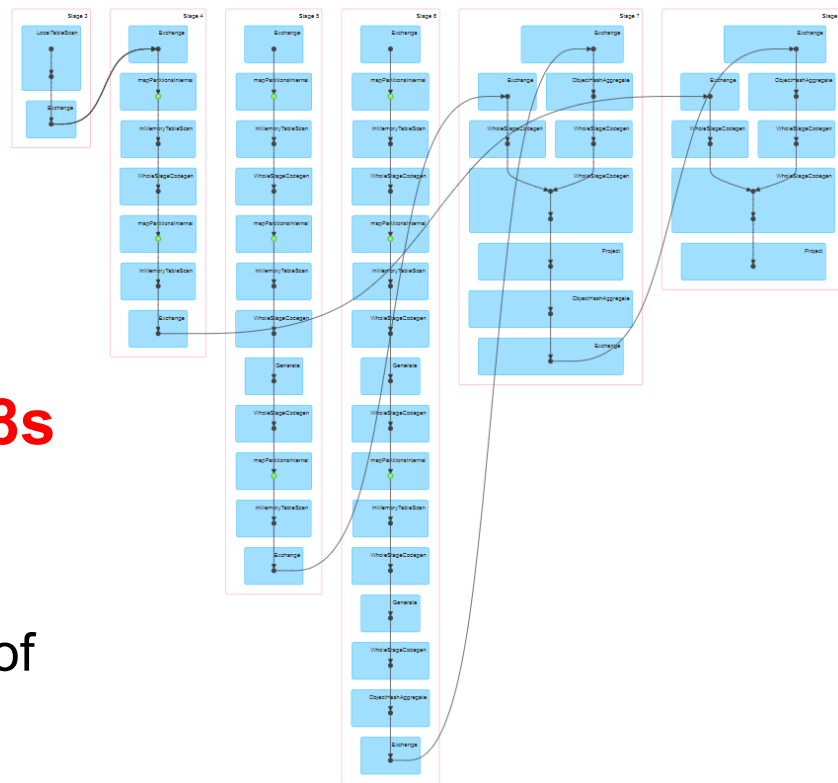
- array_distinct - [SPARK-23912](#)
- array_intersect - [SPARK-23913](#)
- array_union - [SPARK-23914](#)
- array_except - [SPARK-23915](#)
- array_remove - [SPARK-23920](#)
- shuffle - [SPARK-23928](#)
- zip - [SPARK-23931](#)
- map_from_arrays - [SPARK-23933](#)

Higher-order Functions

- transform - [SPARK-23908](#)
- filter - [SPARK-23909](#)
- reduce - [SPARK-23911](#)
- zip_with - [SPARK-23932](#)


Is it really so crucial?

- 1000 rows with 150x150 matrices represented as `array[array[int]]`
 - Square its elements
 - **Explosion approach (47.8s on local[4])**
 - This is a small example
 - What happens with millions of records?



...and with our transform

- With our implementation of transform
 - Still lacks wholestage codegen
 - Very simple implementation
 - Same data, same operation
 - **13.7s on local[4]**



Big Thank to Reviewers!

- **Takuya Ueshin** ~ ueshin
- **Kazuaki Ishizaki** ~ kiskz
- **Liang-Chi Hsieh** ~ viirya
- **Takeshi Yamamuro** ~ maropu
- **Hyukjin Kwon** ~ HyukjinKwon
- **Xiao Li** ~ gatorsmile
- **Wenchen Fan** ~ cloud-fan
- **Weichen Xu** ~ WeichenXu123
- **Marco Gaido** ~ mgaido91
- **Reynold Xin** ~ rxin