





# Image Similarity Detection

## Using LSH and Tensorflow

**Help you discover and do  
what you love.**





100b+  
Pins

2b+  
Boards

People on  
Pinterest  
each month

200m+

10b+  
Recommendations/Day



# Agenda

- 1 Neardup, clustering and LSH**
- 2 Candidate generation**
- 3 Deep dive**
- 4 Candidate selection**
- 5 TF on Spark**

A solid red vertical bar is positioned on the far left side of the image, extending from the top to the bottom.

# Neardup



Artichokes make a bold-back summer appetizer, hosts Greg Denton and Gabriele Guíñón Denton say, along with corn- and chili-topped roasted ricotts.

**FIRE-ROASTED RICOTTA  
WITH CORN AND  
PEPPERS**  
HANDS-ON TIME 10 min

TOTAL TIME 20 MIN.

- 1 recipe Homemade Ricotta, page 136, or one 15-oz. carton whole milk ricotta cheese
- 1 recipe Charred Corn, right
- 1 recipe Grilled Padrón Peppers, right
- 1 Tbsp. chopped Italian parsley
- 1/4 tsp. flaky sea salt
- 1 recipe Balsamic Brown Butter, page 136
- 1 recipe Garlic-Rubbed Bread,

Heat grill to medium-high. Place steaks in an 8- or 9-inch cast-iron skillet. Grill about



Artichokes make a bold-back summer appetizer, hosts Greg Denton and Catherine Gutierrez Denton say, along with corn- and chili-topped roasted meats.

**FIRE-ROASTED RICOTTA  
WITH CORN AND  
PEPPERS**  
HANDS-ON TIME 10 min

total time 20 min.

- recipe Homemade Ricotta,  
page 136, or one 15-oz. carton  
whole milk ricotta cheese  
recipe Cheesed Cann, right  
recipe Grilled Padron Peppers, right  
Tip: chopped Italian parsley  
tip. Sticky sea salt  
recipe Balsamic Brown Butter,  
page 138  
recipe Garlic-Rubbed Bread,  
page 138

Heat grill to medium-high. Place steaks on grill. Cook 5 to 6 minutes on each side. Grill about



# Not Neardup





Unrelated



Neardup



Duplicate



# Clustering

# Not An Equivalence Class

## Formulation

For each image find a canonical image which represents an equivalence class.

## Problem

Neardup is not an equivalence relation because neardup relation is not a transitive relation.

It means we can not find a perfect partition such that all images within a cluster are closer to each other than to the other clusters.



# Incremental approximate K-Cut

Incrementally:

1. **Generate candidates via batch LSH search**
2. **Select candidates via a TF model**
3. Take a transitive closure over selected candidates
4. Pass over clusters and greedily select sub-clusters (K-Cut).



A solid red vertical bar runs along the left edge of the image.

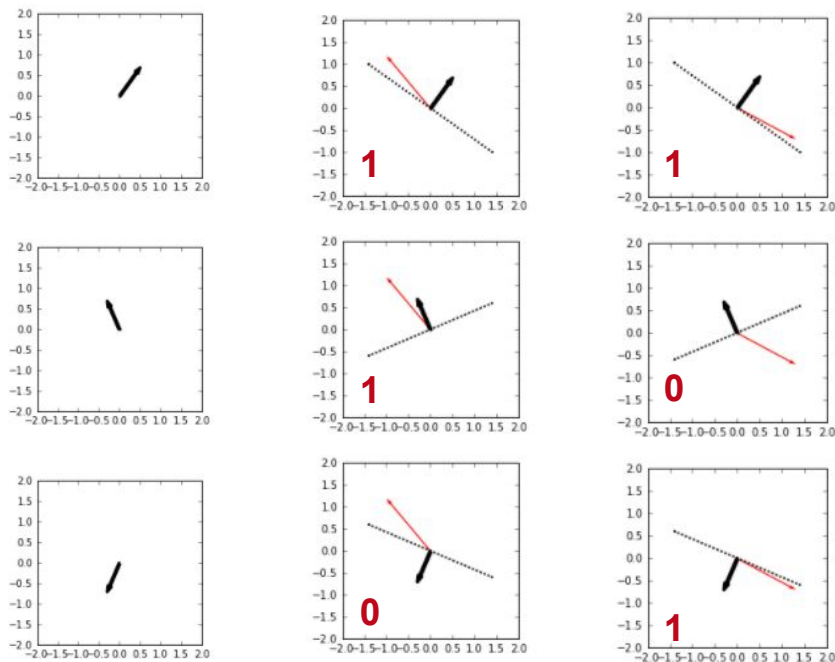
LSH

# Embeddings and LSH

- **Visual Embeddings** are high-dimensional vector representations of entities (in our case images) which capture semantic similarity.
  - Produced via Neural Networks like VGG16, Inception, etc.
- Locality-sensitive hashing or **LSH** is a modern technique used to reduce dimensionality of high-dimensional data while preserving pairwise distances between individual points.



# LSH: Locality Sensitive Hashing



- Pick random projection vectors (**black**)
- For each **embeddings vector** determine on which side of the hyperplane the embeddings vector lands
- On the same side: set bit to 1
- On different side: set bit to 0

Result 1:  $\langle 1 \ 1 \ 0 \rangle$   
Result 2:  $\langle 1 \ 0 \ 1 \rangle$




# LSH terms

Pick optimal number of terms and bits per term

- $1001110001011000 \rightarrow [00]1001 - [01]1100 - [10]0101 - [11]1000$
- $[x] \rightarrow a \text{ term index}$





A solid red vertical bar is positioned on the far left side of the image, extending from the top to the bottom.

# Candidate Generation

# Neardup Candidate Generation

- Input Data:

*RDD[(ImgId, List[LSHTerm])] // billions*

- Goal:

*RDD[(ImgId, TopK[(ImgId, Overlap)])]*

**Nearest Neighbor (KNN) problem formulation**



# Neardup Candidate Generation

Given a set of documents each described by LSH terms, example:

**A**  $\rightarrow$  (1,2,3)

**B**  $\rightarrow$  (1,3,10)

**C**  $\rightarrow$  (2,10)

And more generally:

$D_i \rightarrow [t_j]$

Where each  $D_i$  is a document and  $[t_j]$  is a list of LSH terms (assume each is a 4 byte integer)

Results:

**A**  $\rightarrow$  (B,2) , (C,1)

**B**  $\rightarrow$  (A,2) , (C,1)

**C**  $\rightarrow$  (A,1) , (B,1)



# Spark Candidate Generation

1. Input  $RDD[(ImgId, List[LSHTerm])] \leftarrow$  both **index** and **query** sets
2. flatMap, groupBy input into  $RDD[(LSHTerm, PostingList)] \leftarrow$  an **inverted index**
3. flatMap, groupBy into  $RDD[(LSHTerm, PostingList)] \leftarrow$  a **query list**
4. Join (2) and (3), flatMap over queries posting list, and groupBy query ImgId;  
 $RDD[(ImgId, List[PostingList])] \leftarrow$  **search results by query**.
5. Merge  $List[List[ImgId]]$  into  $TopK(ImgId, Overlap)$  counting number of times each ImgId is seen  $\rightarrow RDD[ImgId, TopK(ImgId, Overlap)]$ .

\*  $PostingList = List[ImgId]$



**Orders of magnitude too slow.**



A solid red vertical bar runs along the left edge of the image.

# Deep Dive

# Dictionary encoding

```
def mapDocToInt(termIndexRaw: RDD[(String, List[TermId])]): RDD[(String, DocId)] = {  
  // ensure that mapping between string and id is stable by sorting  
  // this allows attempts to re-use partial stage completions  
  termIndexRaw.keys.distinct().sortBy(x => x).zipWithIndex()  
}
```

```
val stringArray = (for (ind <- 0 to 1000) yield randomString(32)).toArray  
val intArray = (for (ind <- 0 to 1000) yield ind).toArray
```

108128 Bytes\*  
4024 Bytes\*  
**25x**

\* <https://www.javamex.com/classmexer/>



# Variable Byte Encoding

docIDs	824	829	215406
gaps		5	214577
VB code	00000110 10111000	10000101	00001101 00001100 10110001

- One bit of each byte is a continuation bit; overhead
- int  $\rightarrow$  byte (best case)
- 32 char string up to  $25 \times 4 = \underline{100x}$  memory reduction





# Inverted Index Partitioning

Inverted index is skewed

```
/**
 * Build partitioned inverted index by taking module of docId into partition.
 */
def buildPartitionedInvertedIndex(flatTermIndexAndFreq: RDD[(TermId, (DocId, TermFreq))]):
  RDD[((TermId, TermPartition), Iterable[DocId])] = {

    flatTermIndexAndFreq.map { case (termId, (docId, )) =>
      // partition documents within the same term to improve balance
      // and reduce the posting list length
      ((termId, (Math.abs(docId) % TERM_PARTITIONING).toByte), docId)
    }.groupByKey()
  }
```



# Packing

(Int, Byte) => Long

**Before:**

*Unsorted: **128.77 MB** in 549ms*

*Sort+Limit: 4.41 KB in **7511ms***

**After:**

*Unsorted: **38.83 MB** in 219ms*

*Sort+Limit: 4.41 KB in **467ms***

```
def packDocIdAndByteIntoLong(docId: DocId, docFreq: DocFreq): Long = {  
  (docFreq.toLong << 32) | (docId & 0xffffffffL)  
}  
  
def unpackDocIdAndByteFromLong(packed: Long): (DocId, DocFreq) = {  
  (packed.toInt, (packed >> 32).toByte)  
}
```



# Slicing

Split query set into slices to reduce spill and size for “widest” portion of the computation. Union at the end.



# Additional Optimizations

- **Cost based optimizer** - significant improvements to runtime can be realized by analyzing input data sets and setting performance parameters automatically.
- **Counting** - jaccard overlap counting is done via low level, high performance collections.
- **Off heaping** serialization when possible (*spark.kryo.unsafe*).



# Generic Batch LSH Search

- Can be applied generically to KNN, embedding agnostic.
- Can work on arbitrary large query set via slicing.



A solid red vertical bar is positioned on the far left side of the image, extending from the top to the bottom.

# Candidate Selection

# TF DNN Classifier

- Transfer learning over VGG16
- Visual embeddings
- XOR hamming bits
- Learning still happens at >1B pairs
- Batch size of 1024, Adam optimizer



# Vectorization: mapPartitions + grouped

- During training and inference vectorization reduces overhead.
- Spark mapPartitions + grouped allows for large batches and controlling the size. Works well for inference.
- 2ms/prediction on c3.8xl CPUs with network of 10MM parameters .

```
input.mapPartitions { partition: Iterator[(ImgInfo, ImgInfo)] =>  
  
  // break down large partitions into groups and score per group  
  partition.grouped(BATCH_SIZE).flatMap { group: Seq[(ImgInfo, ImgInfo)] =>  
    // create tensors and score as features: Array[Array[Float]] --> Tensor.create(features)  
    }  
}
```





# One TF Session per JVM

- Reduce model loading overhead, load once per JVM; thread-safe.

```
object TensorflowModel {  
  lazy val model: Session = {  
    SavedModelBundle.load(...).session()  
  }  
}
```



# Summary

- Candidate Generation uses Batch LSH Search over terms from visual embeddings.
- Batch LSH scales to billions of objects in the index and is embedding agnostic.
- Candidate Selection uses a TF classifier over raw visual embeddings.
- Two-pass transitive closure to cluster results.



A solid red vertical bar is positioned on the far left side of the image, extending from the top to the bottom.

**Thanks!**